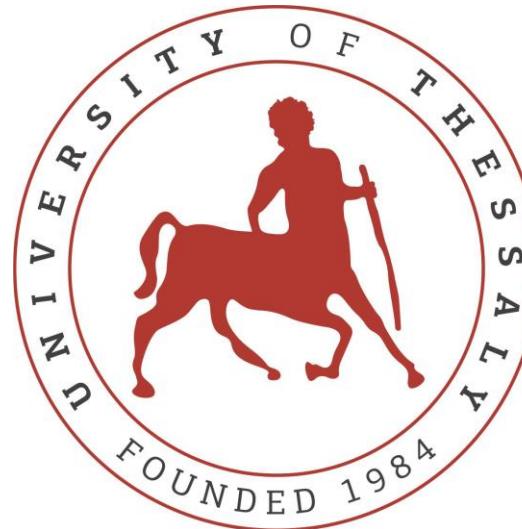


---

# ECE494 MICROPROCESSOR DESIGN

## PIPELINED FFT IMPLEMENTATION



Christodoulos Zerdalis 03531

Athanasiос Gerampinis 03466

26/6

---

# UNDERSTANDING THE FAST FOURIER TRANSFORM

- *What is FFT?*
    - An optimized algorithm to compute the **Discrete Fourier Transform** (DFT) and its inverse, reducing time complexity from  $O(N^2)$  to  $O(N \log N)$ .
  - *Why is FFT Important?*
    - Essential in **digital signal processing** (DSP)
    - Used in **image processing, audio analysis, wireless communications.**
  - *Basic Principle:* FFT breaks down a **large DFT** into **smaller DFTs** using divide-and-conquer strategies (e.g., Radix-2, Radix-4).
-

---

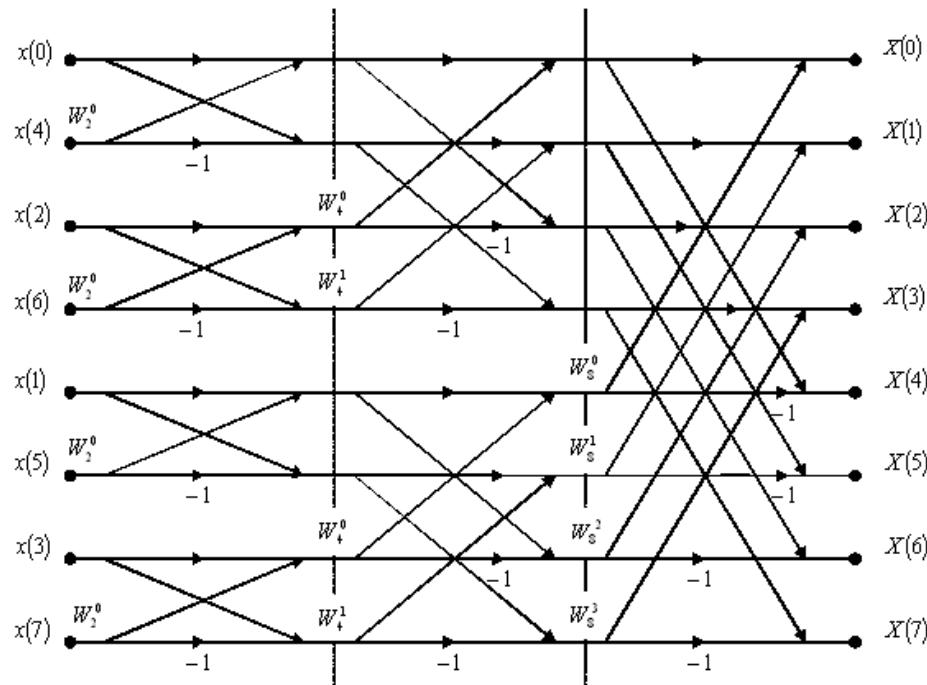
# MOTIVATION FOR FFT ACCELERATION

- **High Performance Required in Real-Time Systems:**
  - Voice recognition, and real-time video processing need fast FFT computation.
- **Energy & Area Constraints in Embedded Systems:**
  - Optimized FFT cores reduce power and hardware resources.
- **Large-Scale Data Processing:**
  - FFT is a core computational kernel in large systems like wireless communication stations

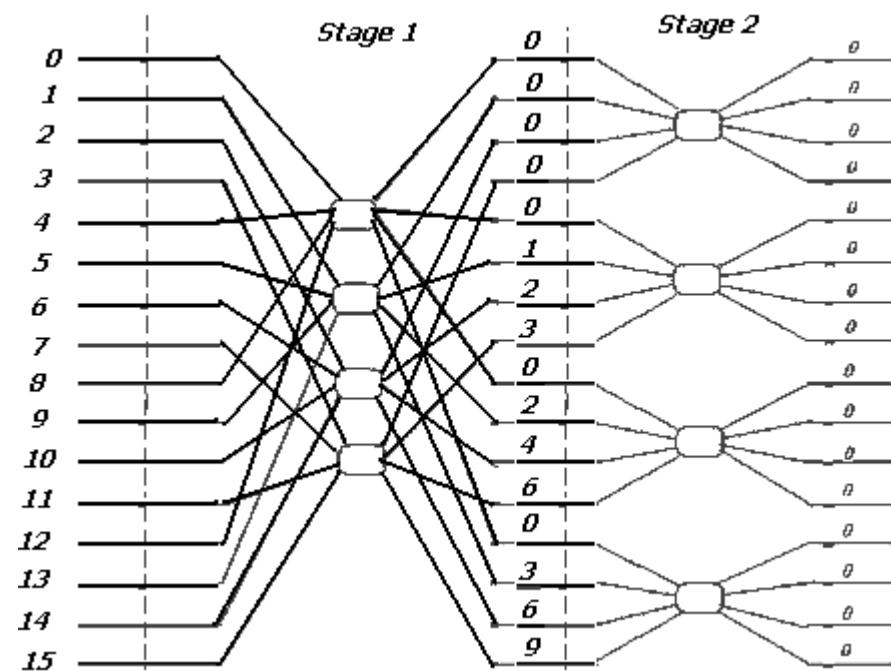
# PROJECT FOCUS

Our goal is to develop hardware implementations of two different FFT algorithms (Radix-2 and Radix-4) and compare them in terms of speed, power consumption, and area for various input sizes.

**Radix-2**



**Radix-4**



# TWIDDLE COMPUTATION

- Since the twiddle factors W involve complex number calculations, their values are **not computed on-the-fly in hardware**. Instead, these values are **precomputed beforehand** and stored in a **lookup table (LUT)**.

*This approach reduces computational complexity and saves hardware resources*

*We save the real and imaginary values in 16-bit buffers using Q1.15 fixed point arithmetic*

Index	W_Real	W_Img	Complex number
0	32768	0	$1 + 0j$
1	23170	-23170	$0.7071 - 0.7071j$
2	0	-32768	$0 - 1j$
3	-23170	-23170	$-0.7071 - 0.7071j$
4	-32768	0	$-1 + 0j$
5	-23170	23170	$-0.7071 + 0.7071j$
6	0	32768	$0 + 1j$
7	23170	23170	$0.7071 + 0.7071j$

---

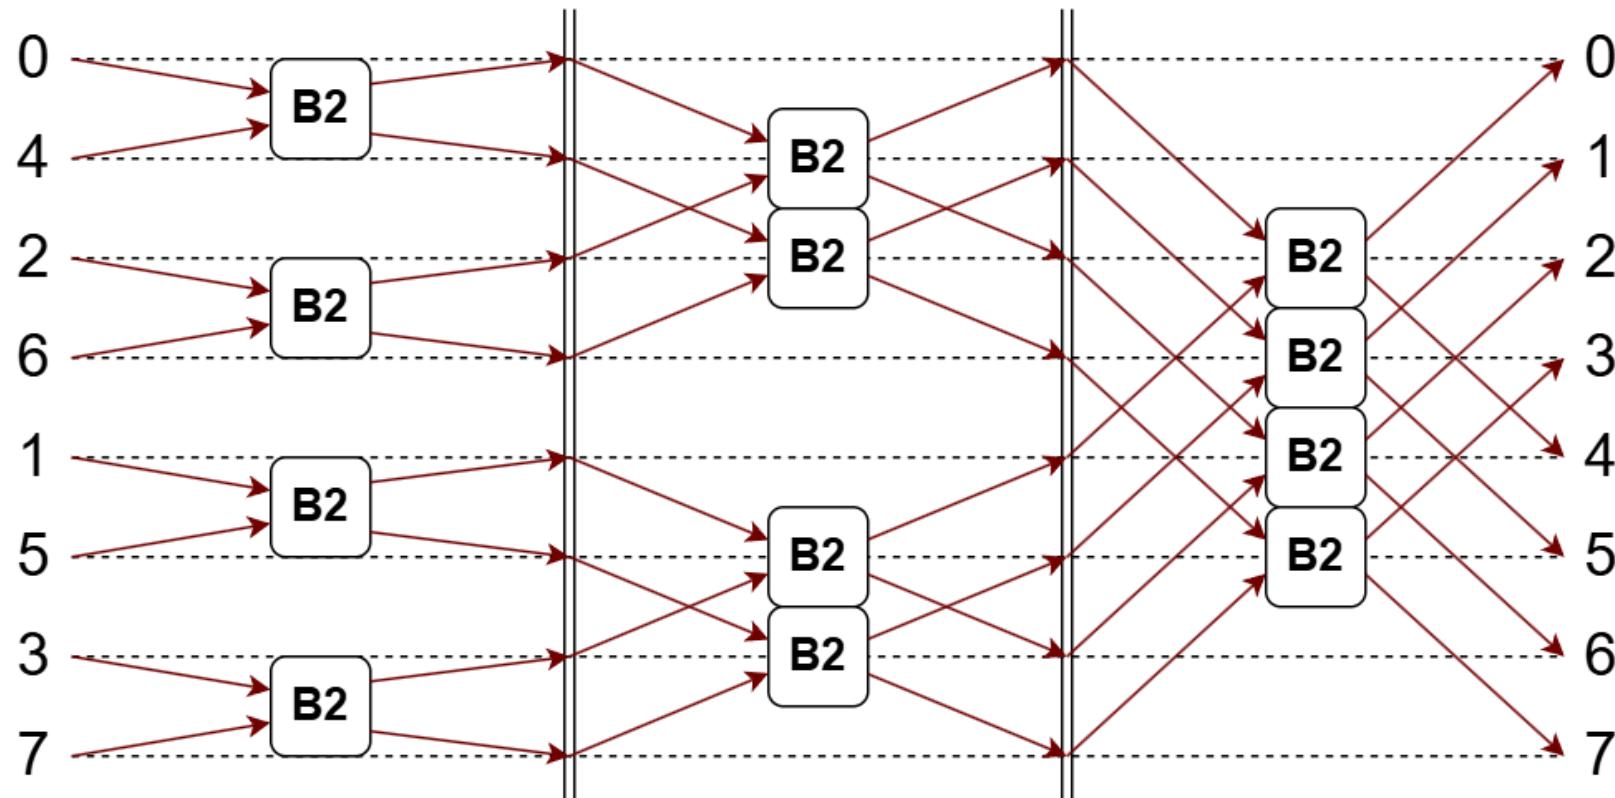
# REORDER INPUT DATA

- Before starting butterfly calculations in an FFT we need to **reorder** the data, for this we use **bit-reversal**

Index	Binary	Bit-Reversed Binary	Bit-Reversed Index
0	000	000	0
1	001	100	4
2	010	010	2
3	011	110	6
4	100	001	1
5	101	101	5
6	110	011	3
7	111	111	7

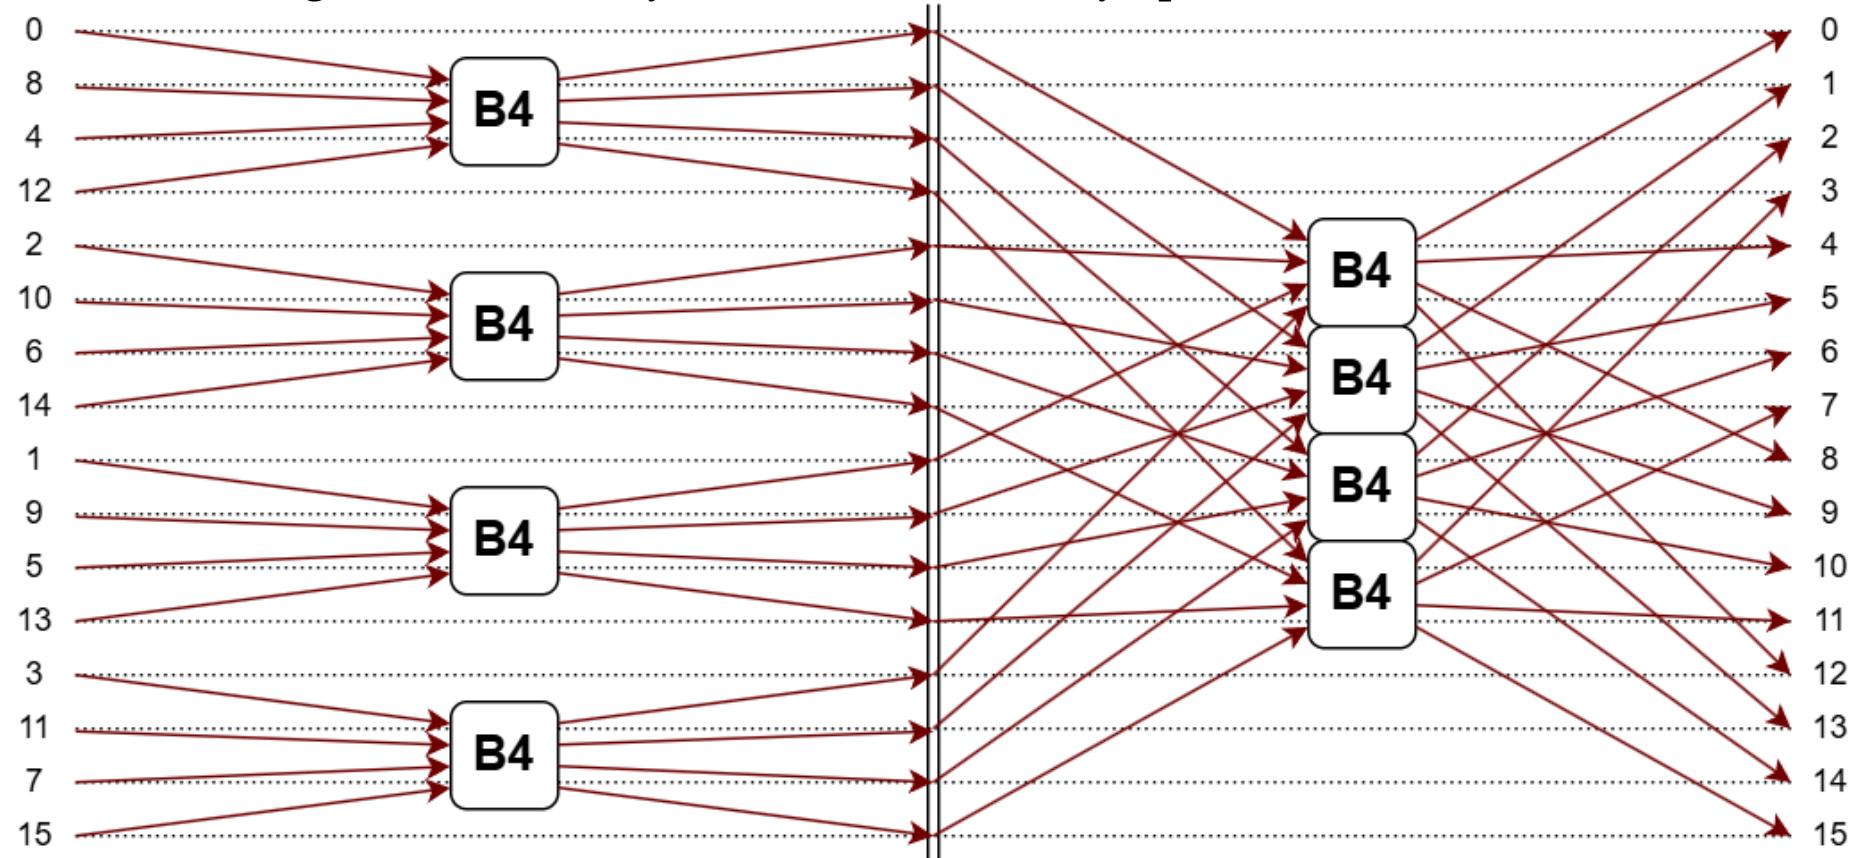
# RADIX-2 (unrolled)

This implementation of the Radix-2 algorithm fully utilizes its parallelism, resulting in a fully unrolled design that uses a dedicated butterfly unit for each butterfly operation.



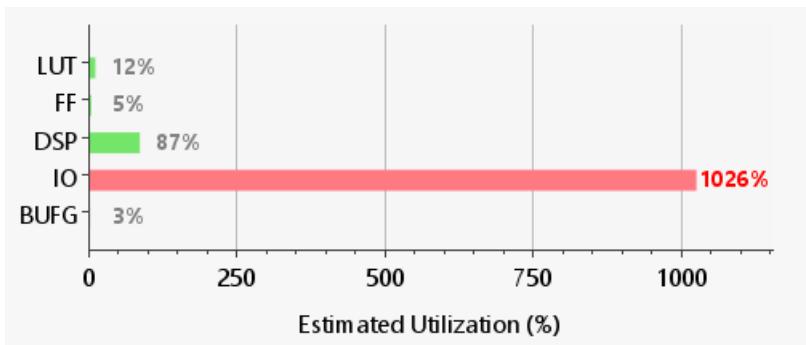
# RADIX-4 (unrolled)

Again, the unrolled design uses a butterfly unit for each butterfly operation.

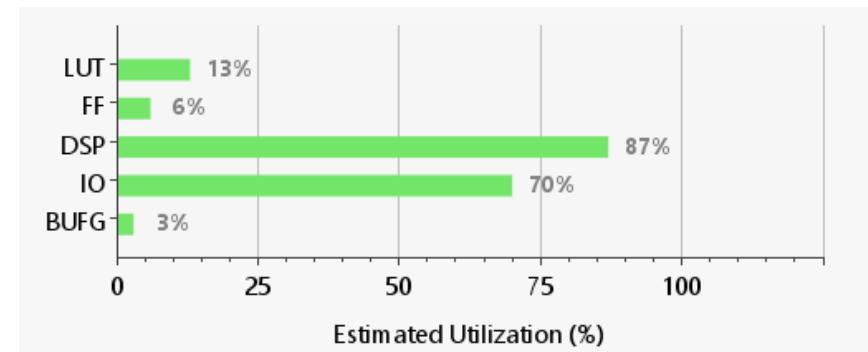


# DATA BANDWIDTH LIMITATION

- These designs require all input **data to be available simultaneously** to begin computation. For example,  $32 \text{ bits} \times 16 \text{ inputs} \times 2 \text{ types} \times 2 \text{ input/output} = 2048 \text{ bits}$ , which far exceeds the **I/O capabilities** of our FPGA. To overcome this, we designed a **memory controller** that inputs data one at a time, stores them until needed, and then outputs the results sequentially.
- This necessary change makes the design slower. Initially, the computation time was  $T = \text{Number\_of\_stages}$ , where  $\text{Number\_of\_stages} = \log(\text{Input\_size})$ . For example,  $T_{16} = 4$  for a 16-point FFT. Now this time is  $T_{\text{new}} == \log(\text{Input\_size}) + 2 \times \text{Input\_size}$ .

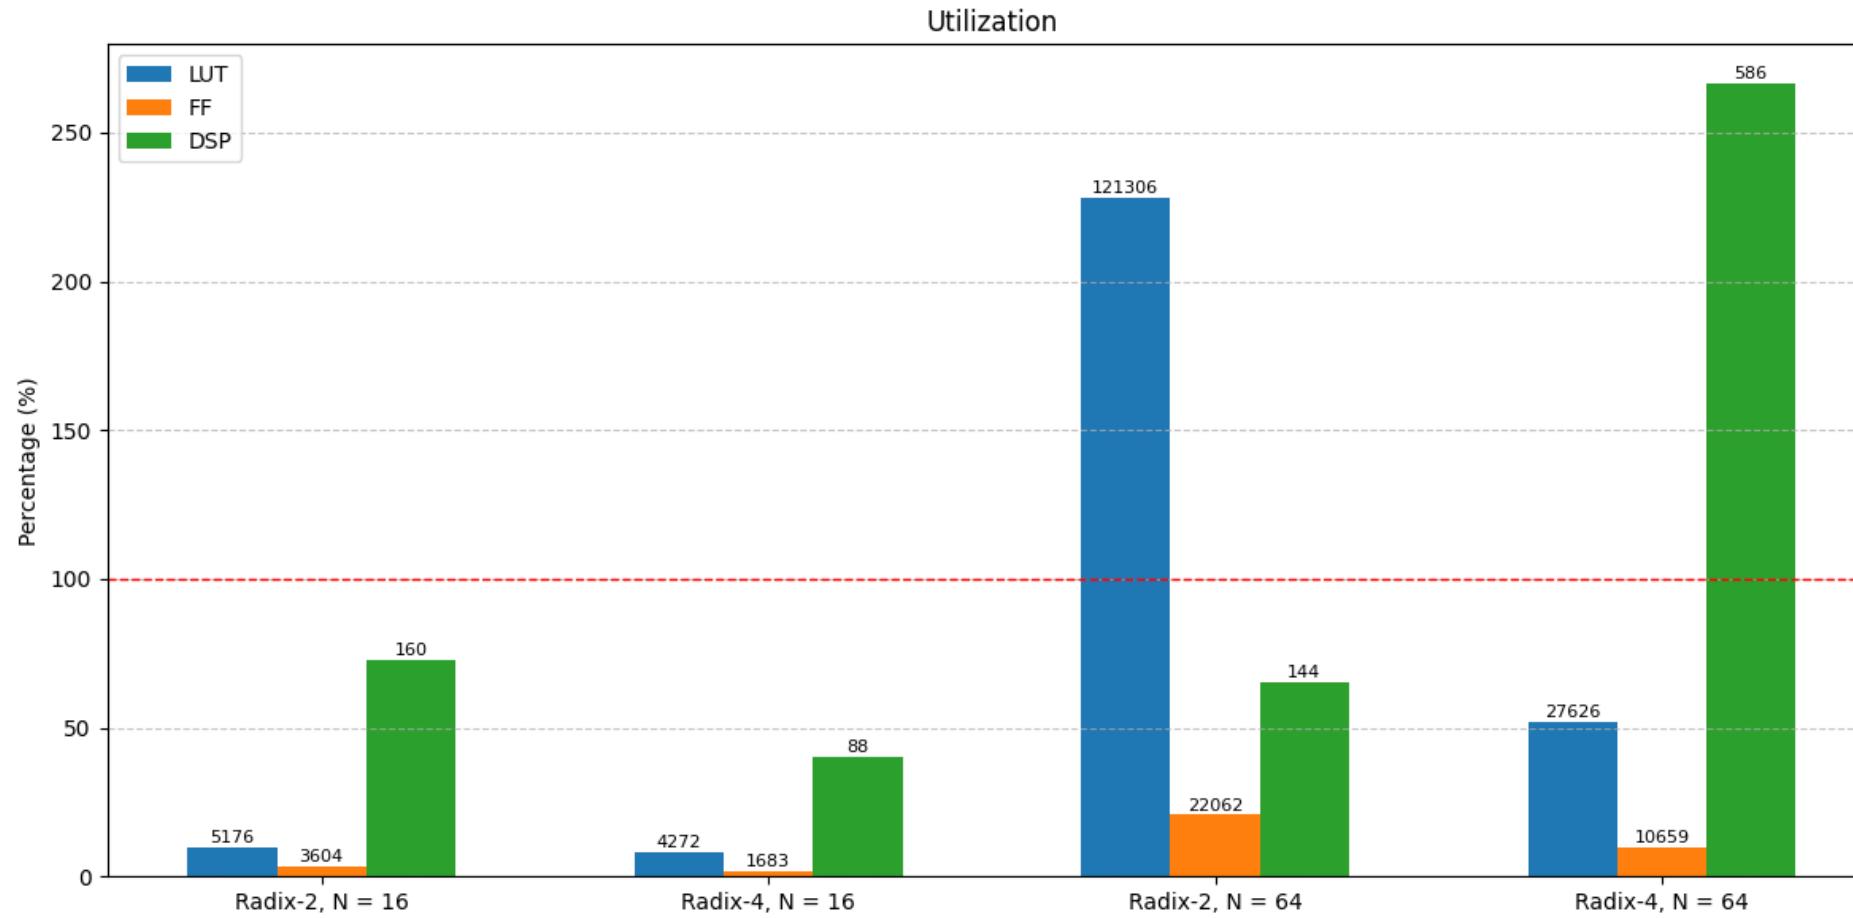


No memory control

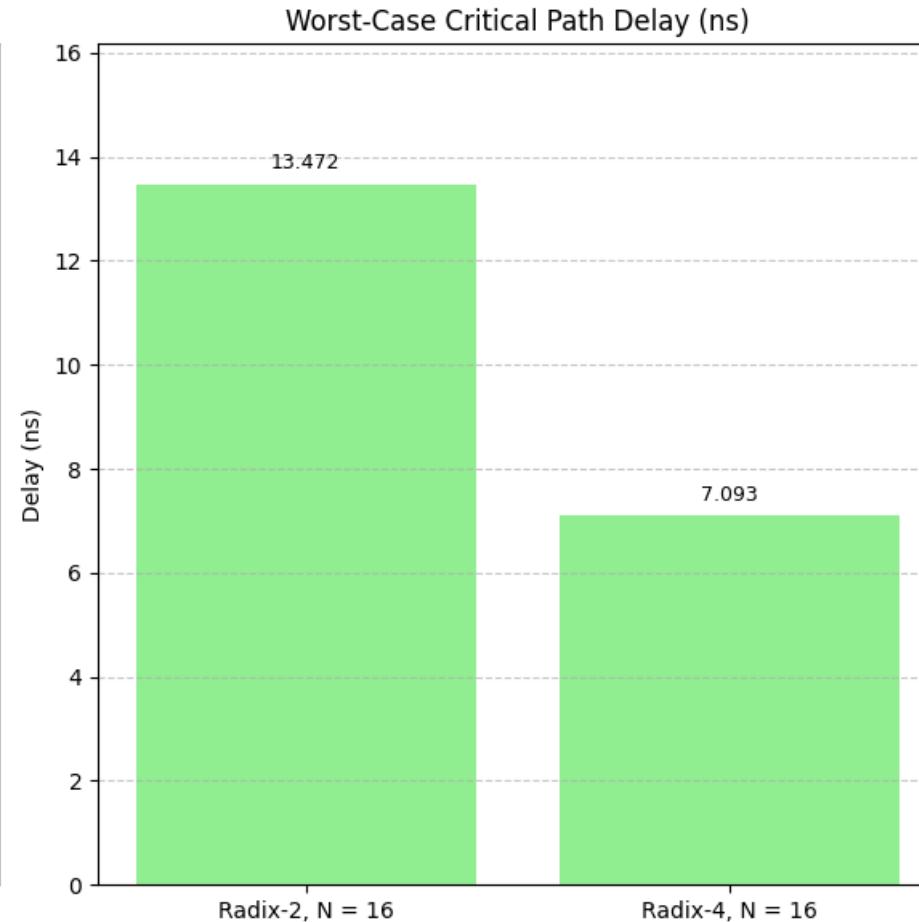
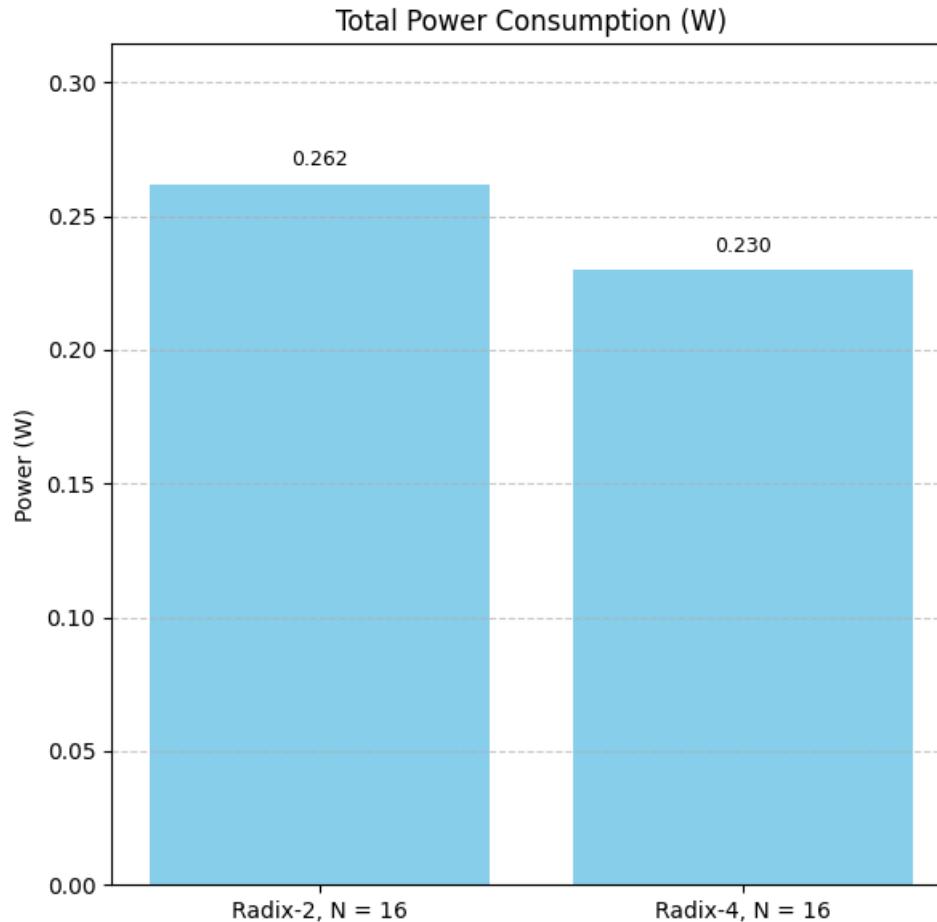


Memory control

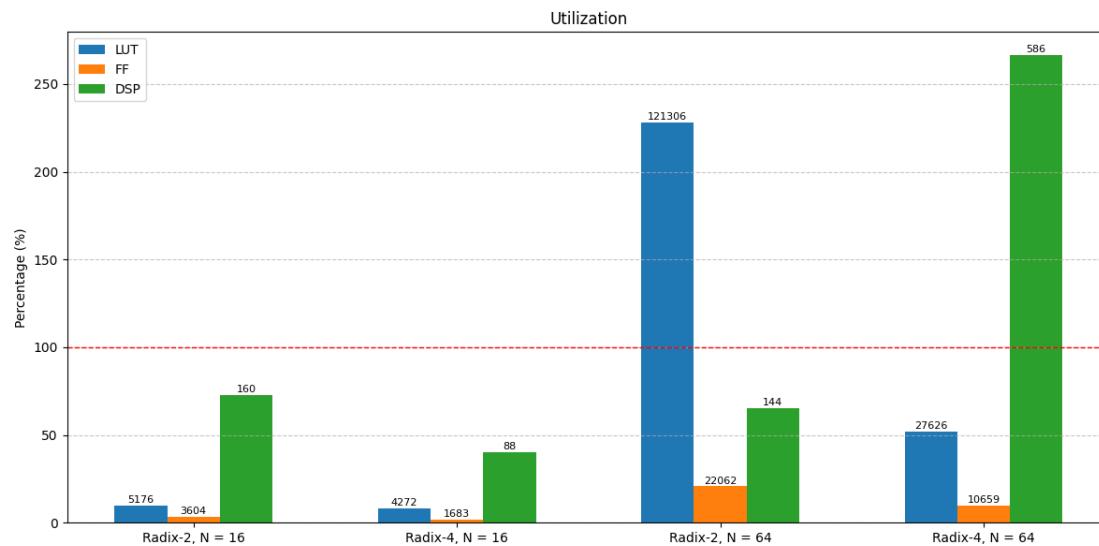
# RESULTS



# RESULTS



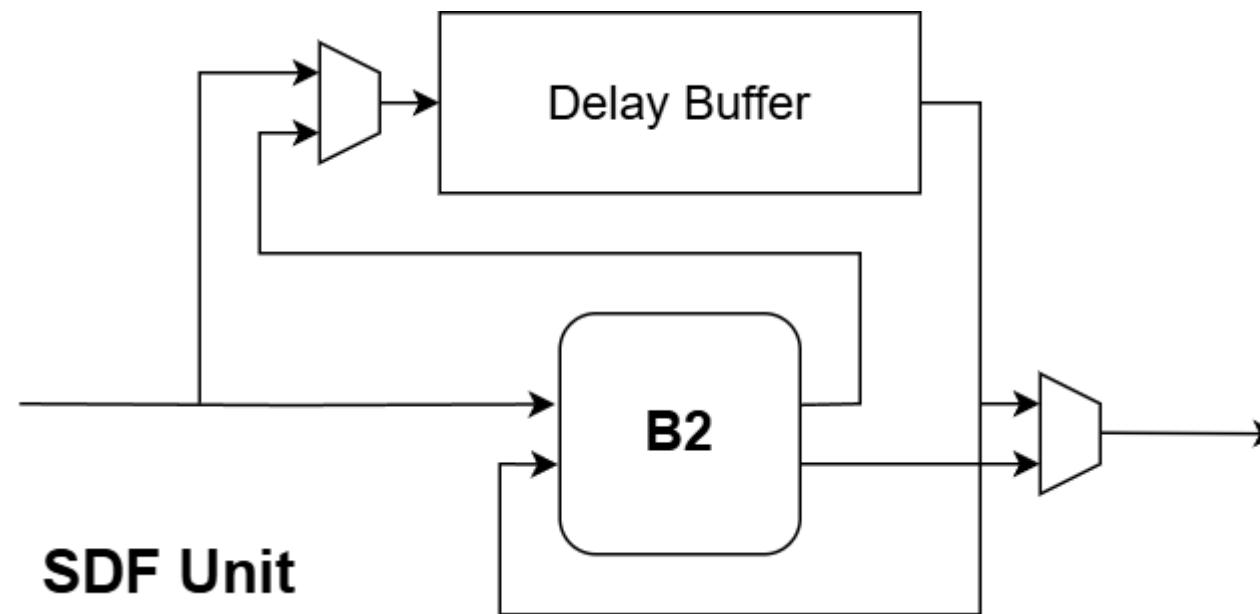
# LIMITATIONS (scalability)



- We can see that as the **input size increases**, the number of **resources and the area required grow significantly**. This becomes a problem—even for a small input size of 64, the design **cannot be implemented** on our FPGA.

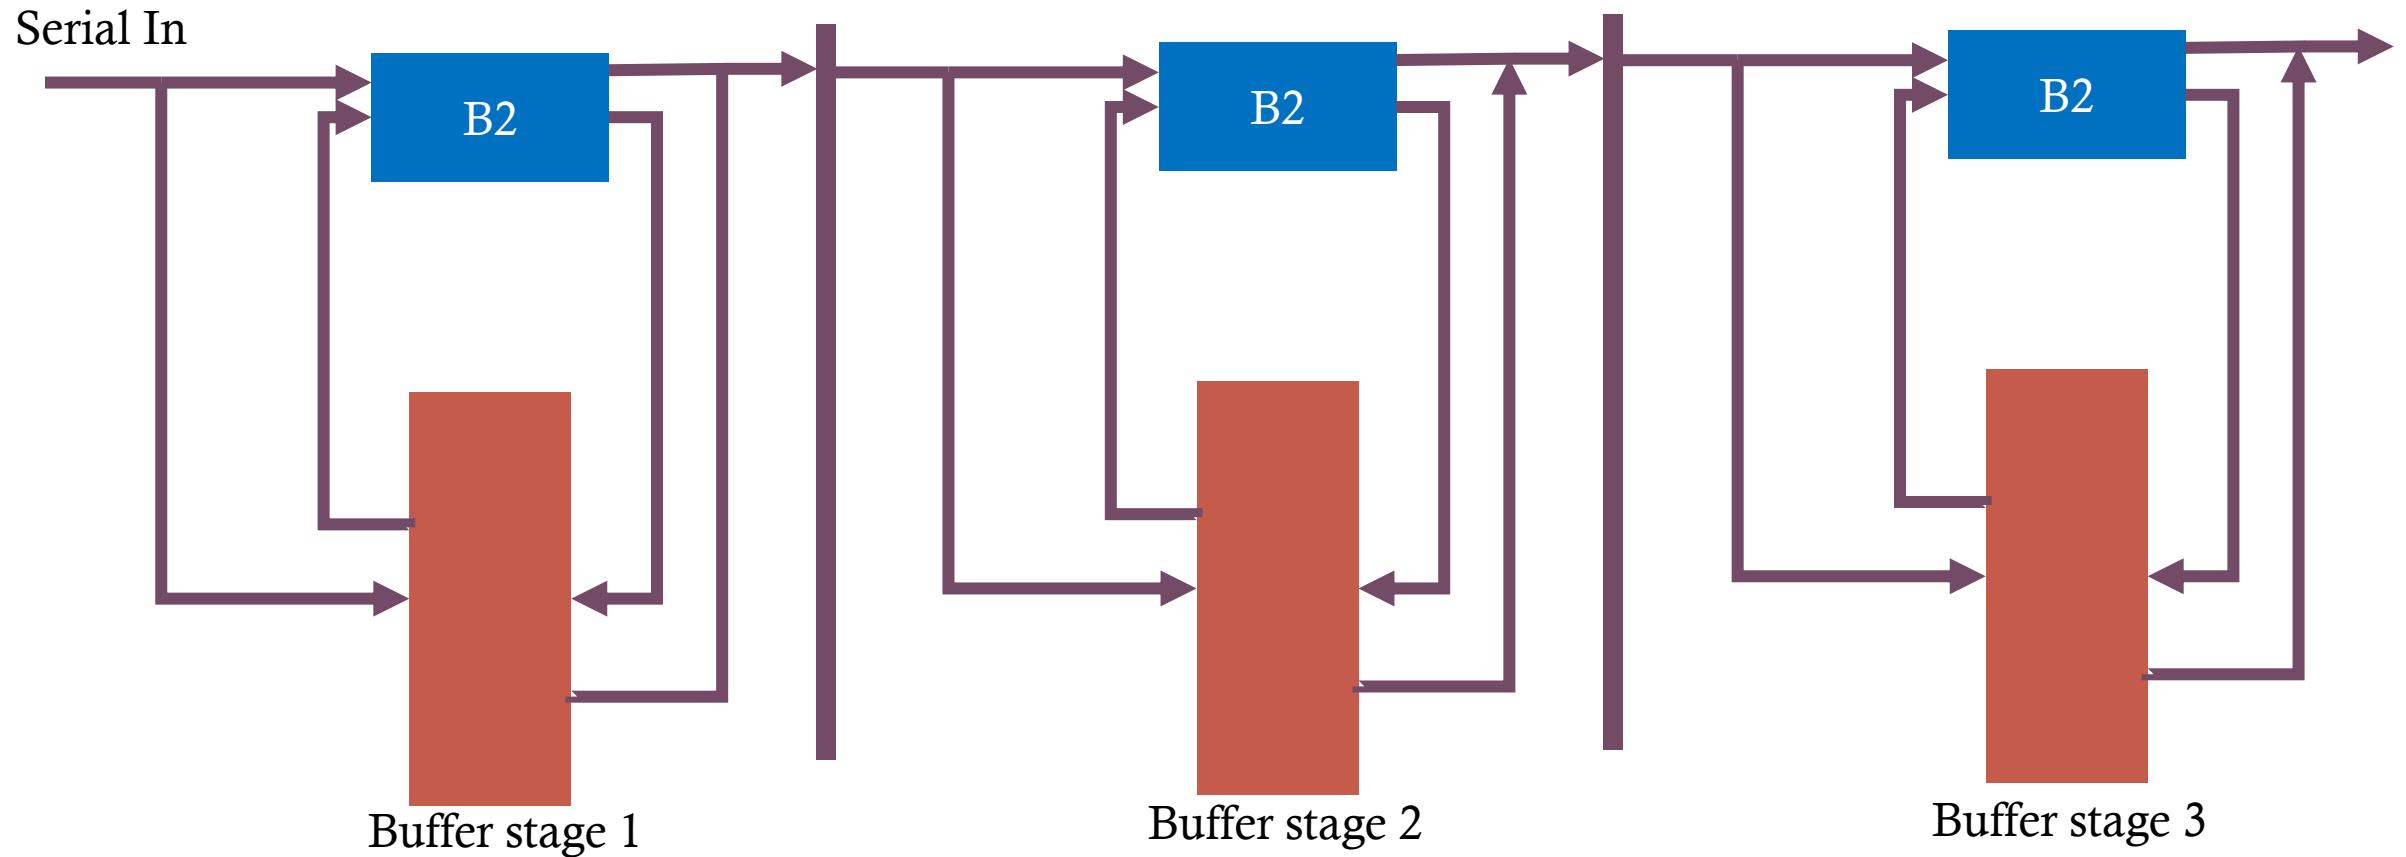
# NEW APPROACH

- Radix-2 SDF performs one stage of radix-2 FFT using a **Single-Path Delay Feedback** architecture by **streaming** complex input samples, applying butterfly and twiddle operations with proper data alignment via a **delay buffer**, and outputting the transformed results **sequentially**.



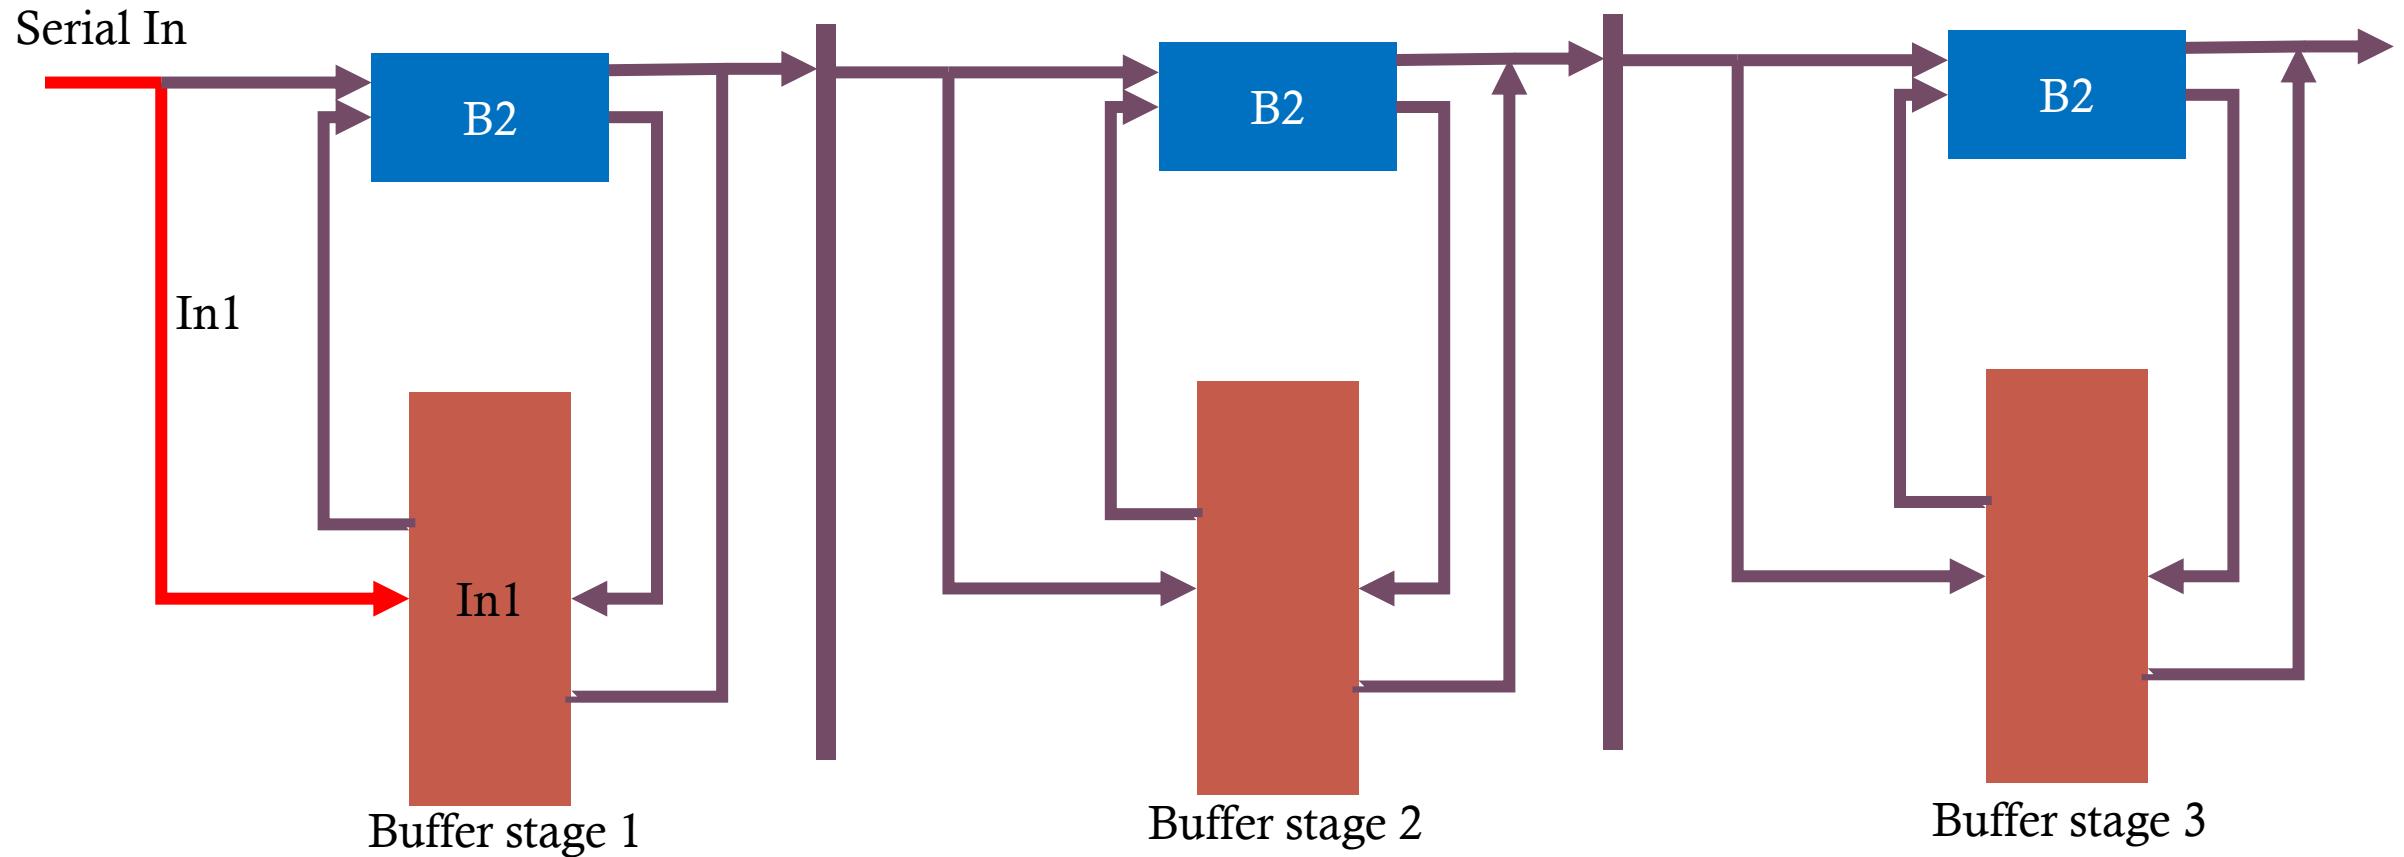
# NEW APPROACH

- Radix-2 SDF (Single path delay feedback)



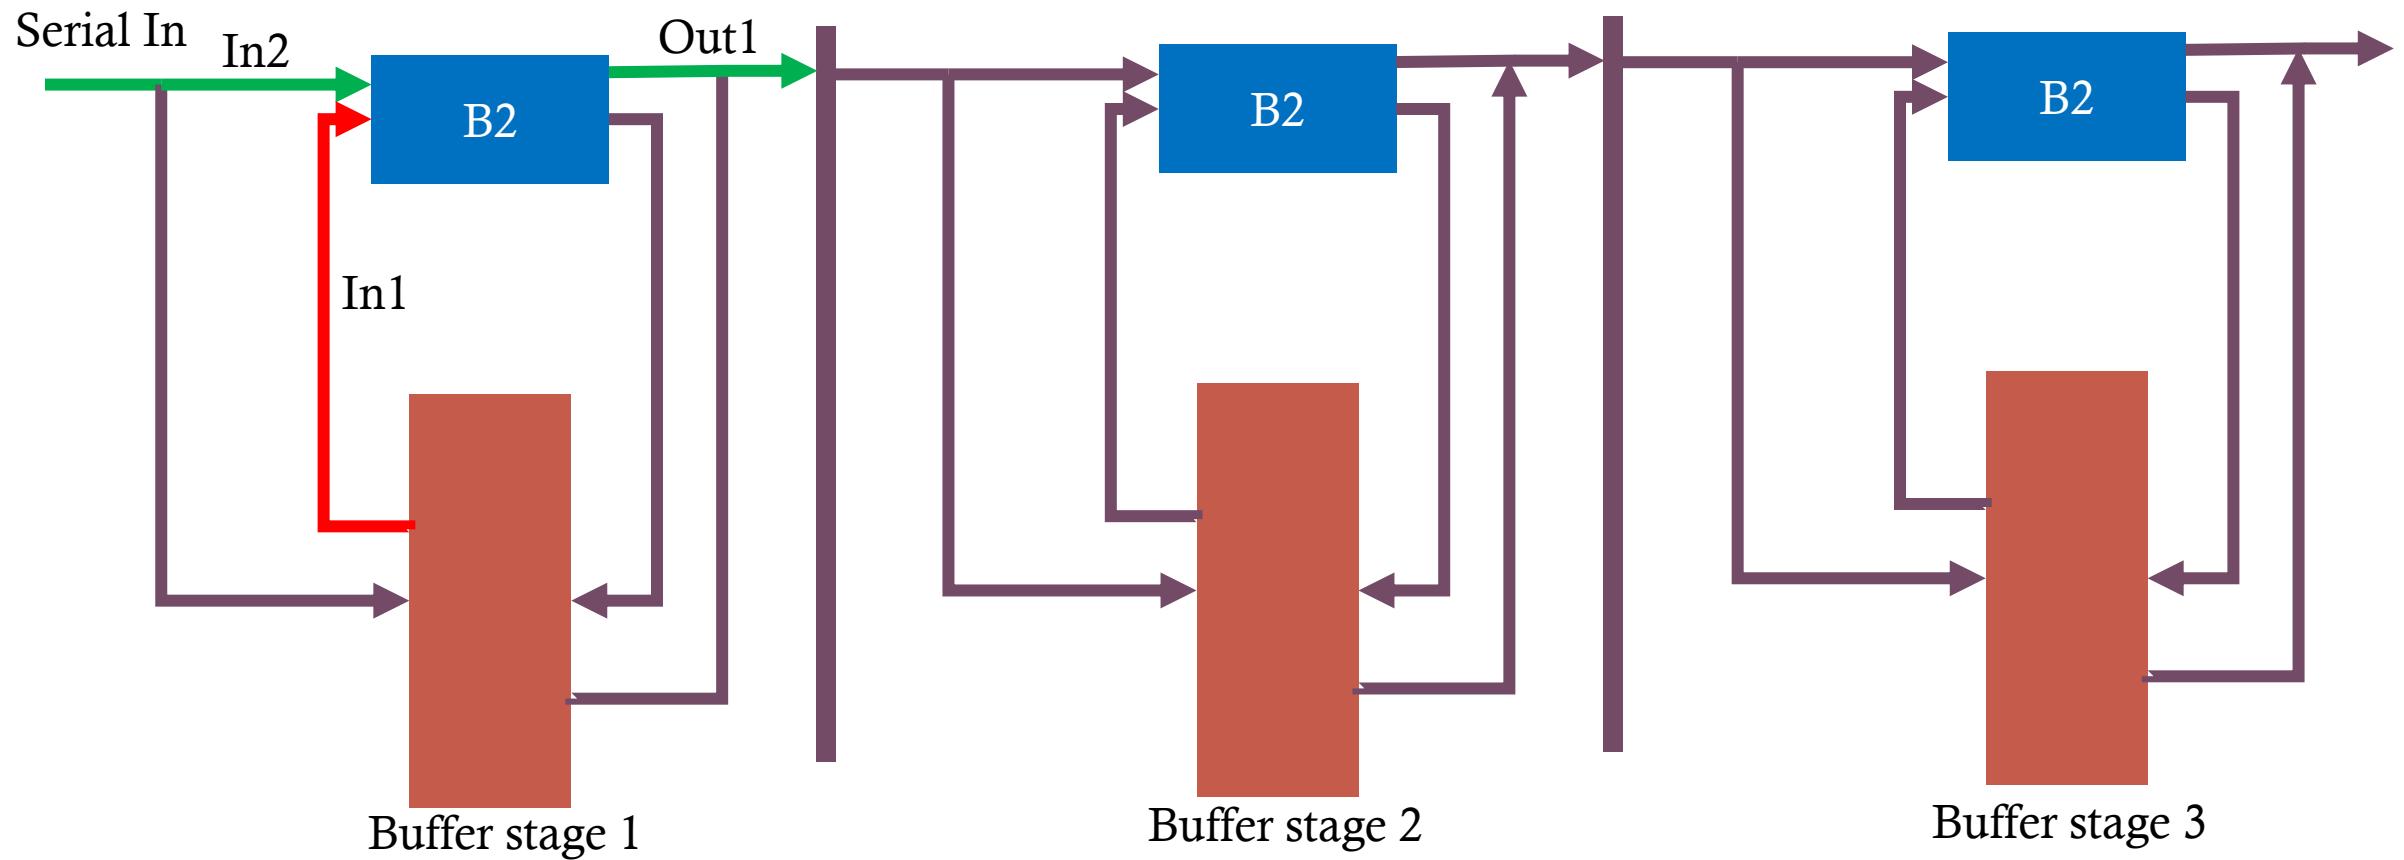
# NEW APPROACH

- Radix-2 SDF (Single path delay feedback)



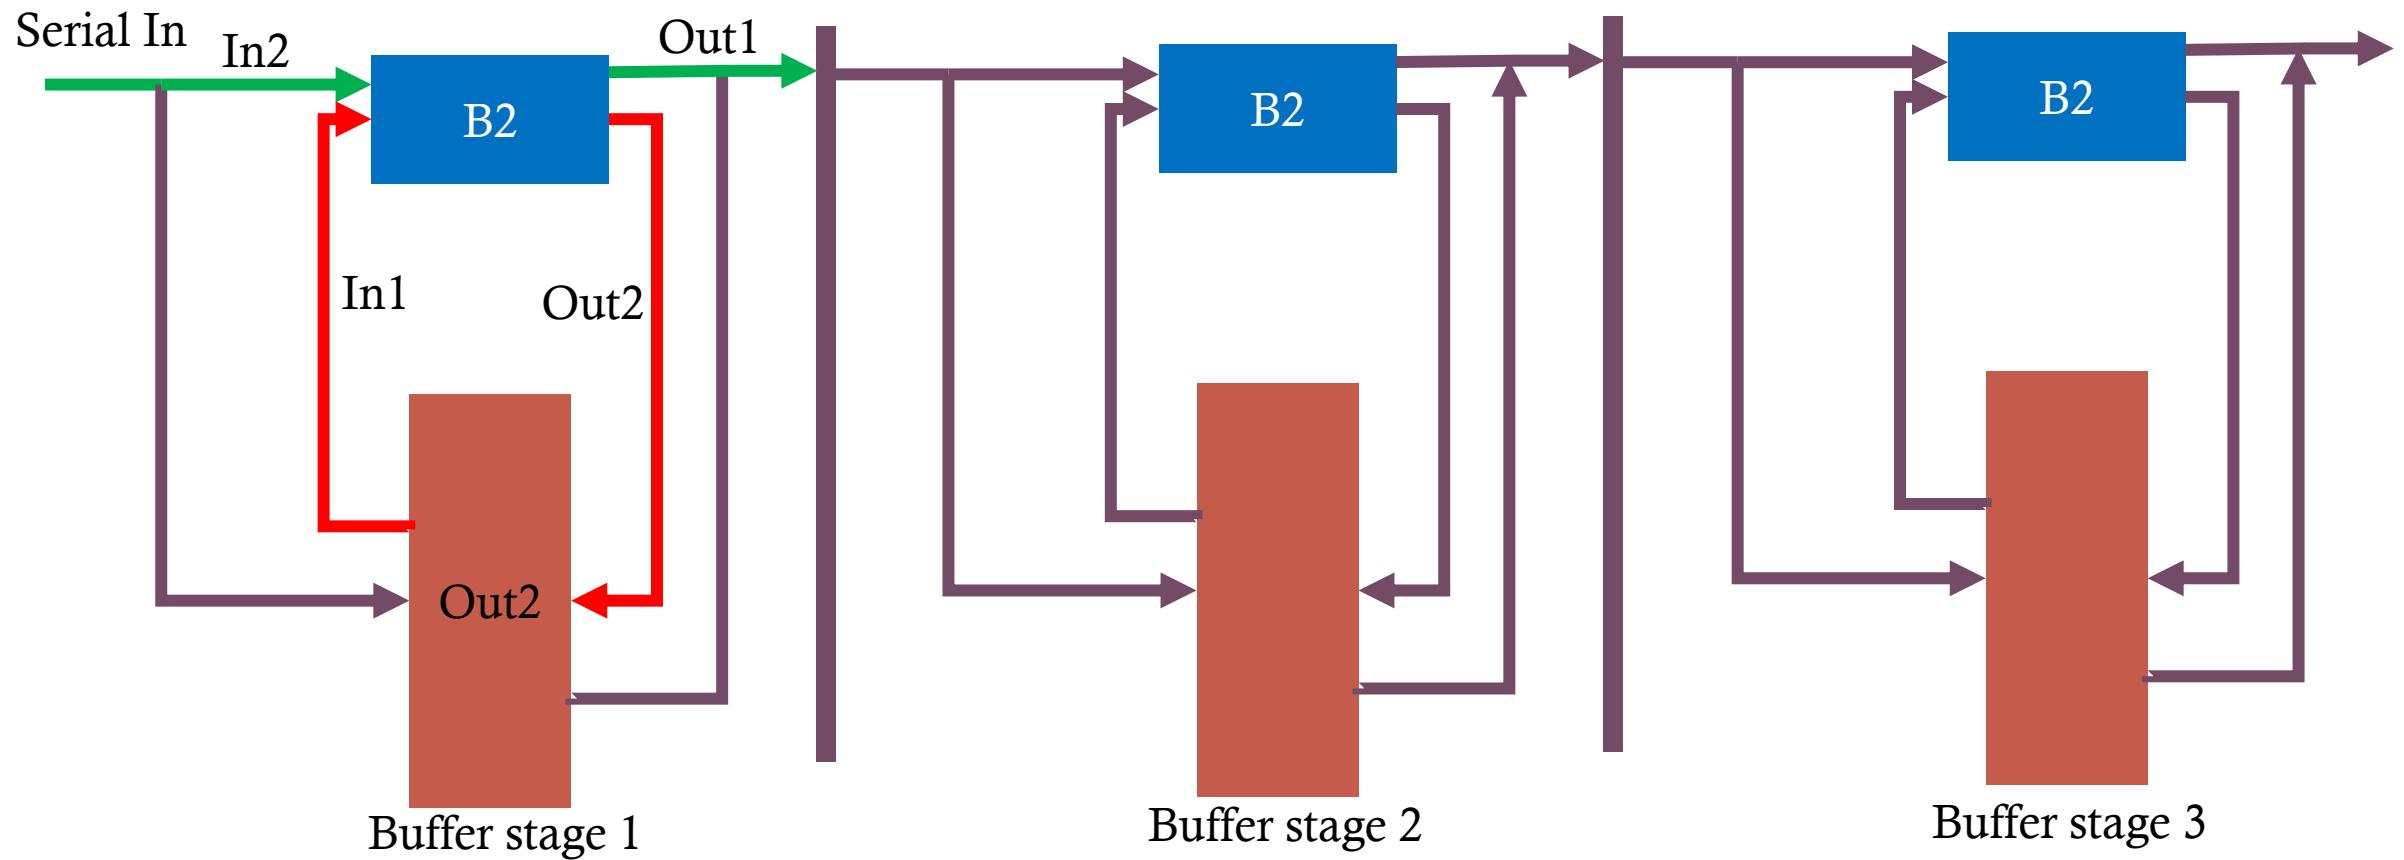
# NEW APPROACH

- Radix-2 SDF (Single path delay feedback)



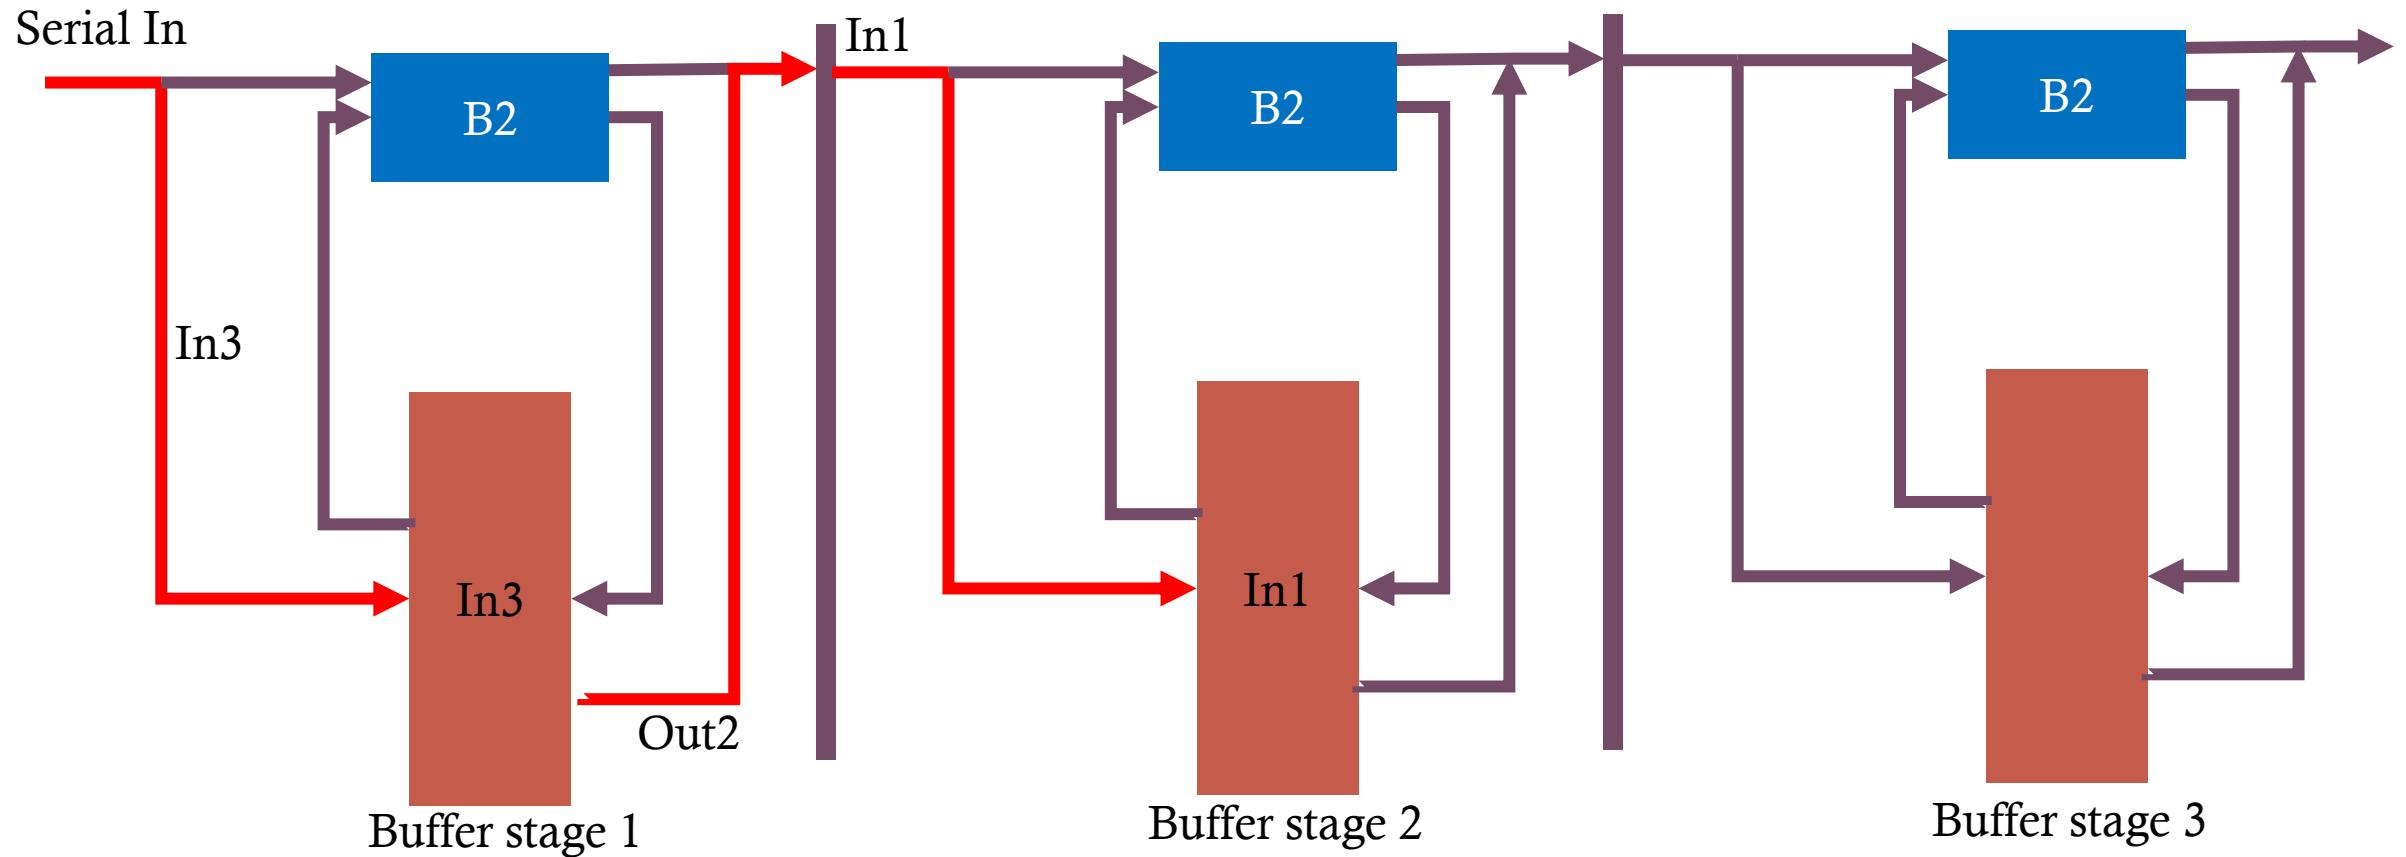
# NEW APPROACH

- Radix-2 SDF (Single path delay feedback)



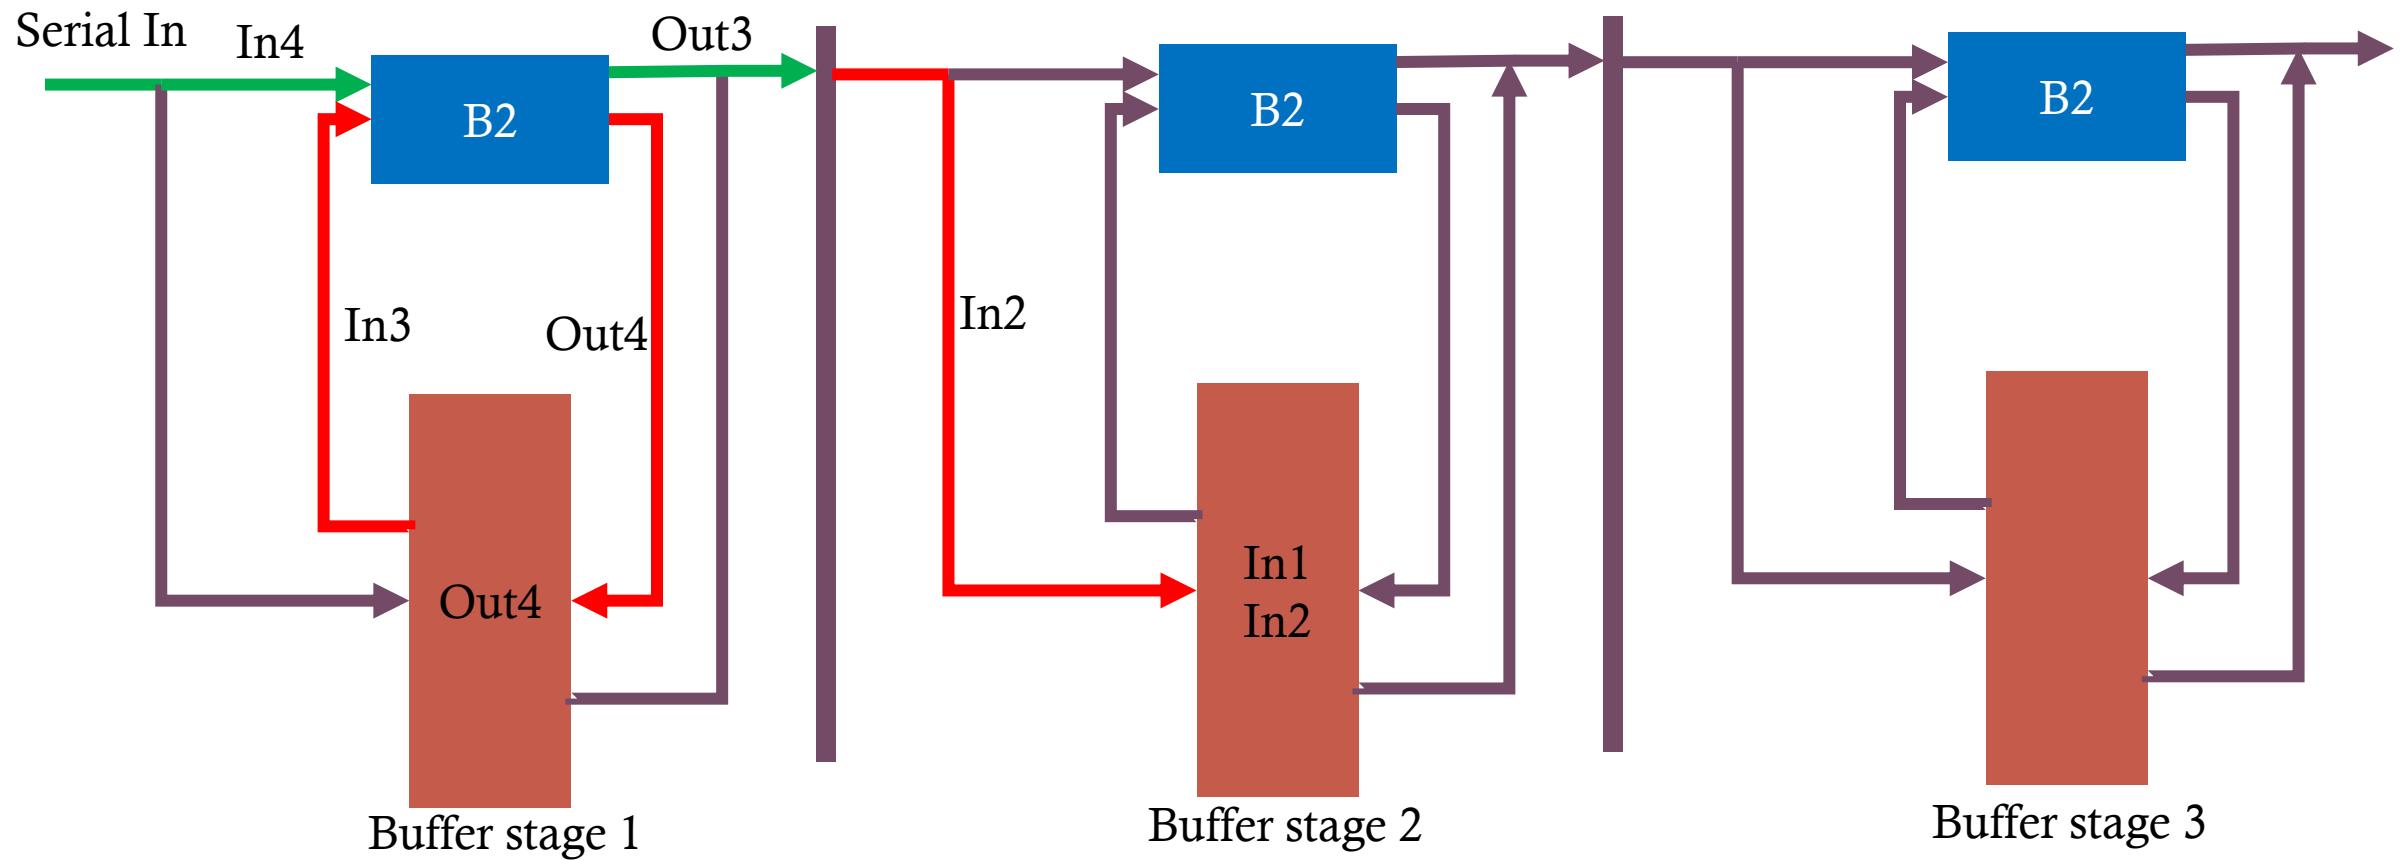
# NEW APPROACH

- Radix-2 SDF (Single path delay feedback)



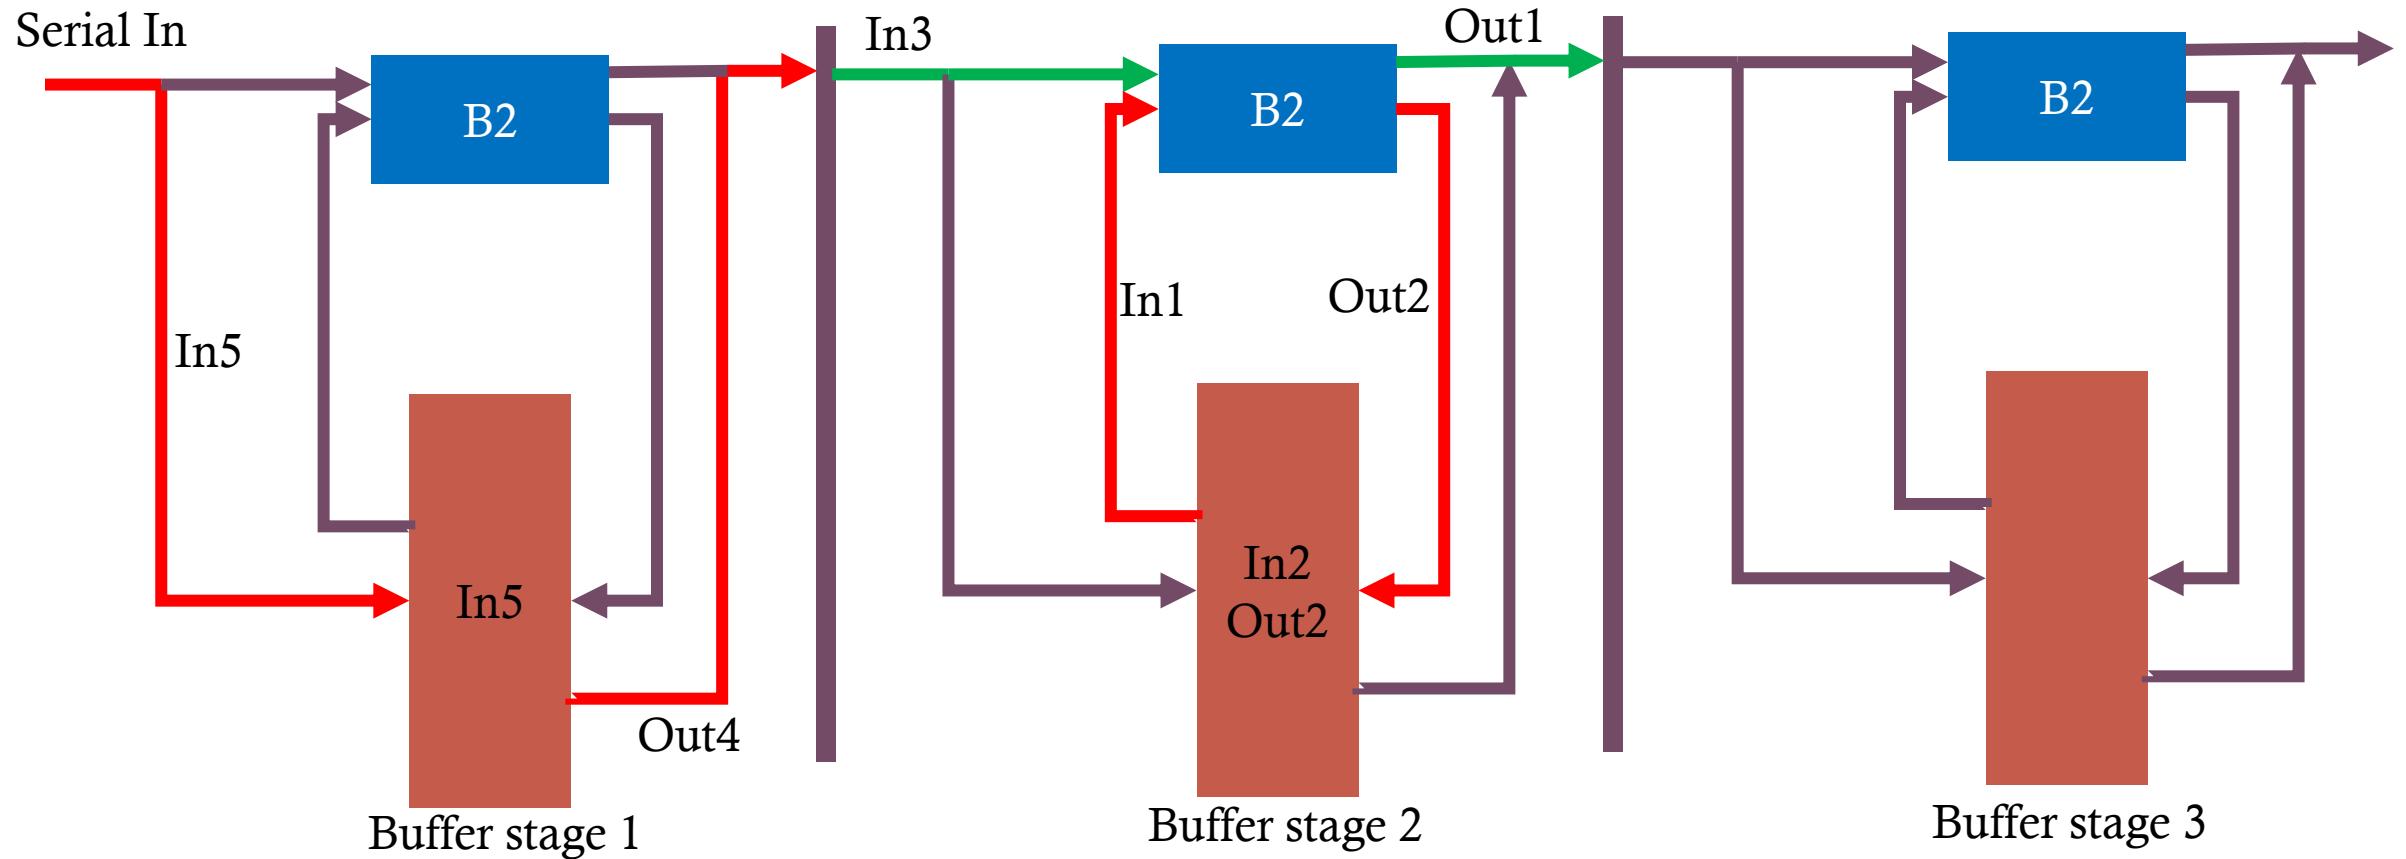
# NEW APPROACH

- Radix-2 SDF (Single path delay feedback)

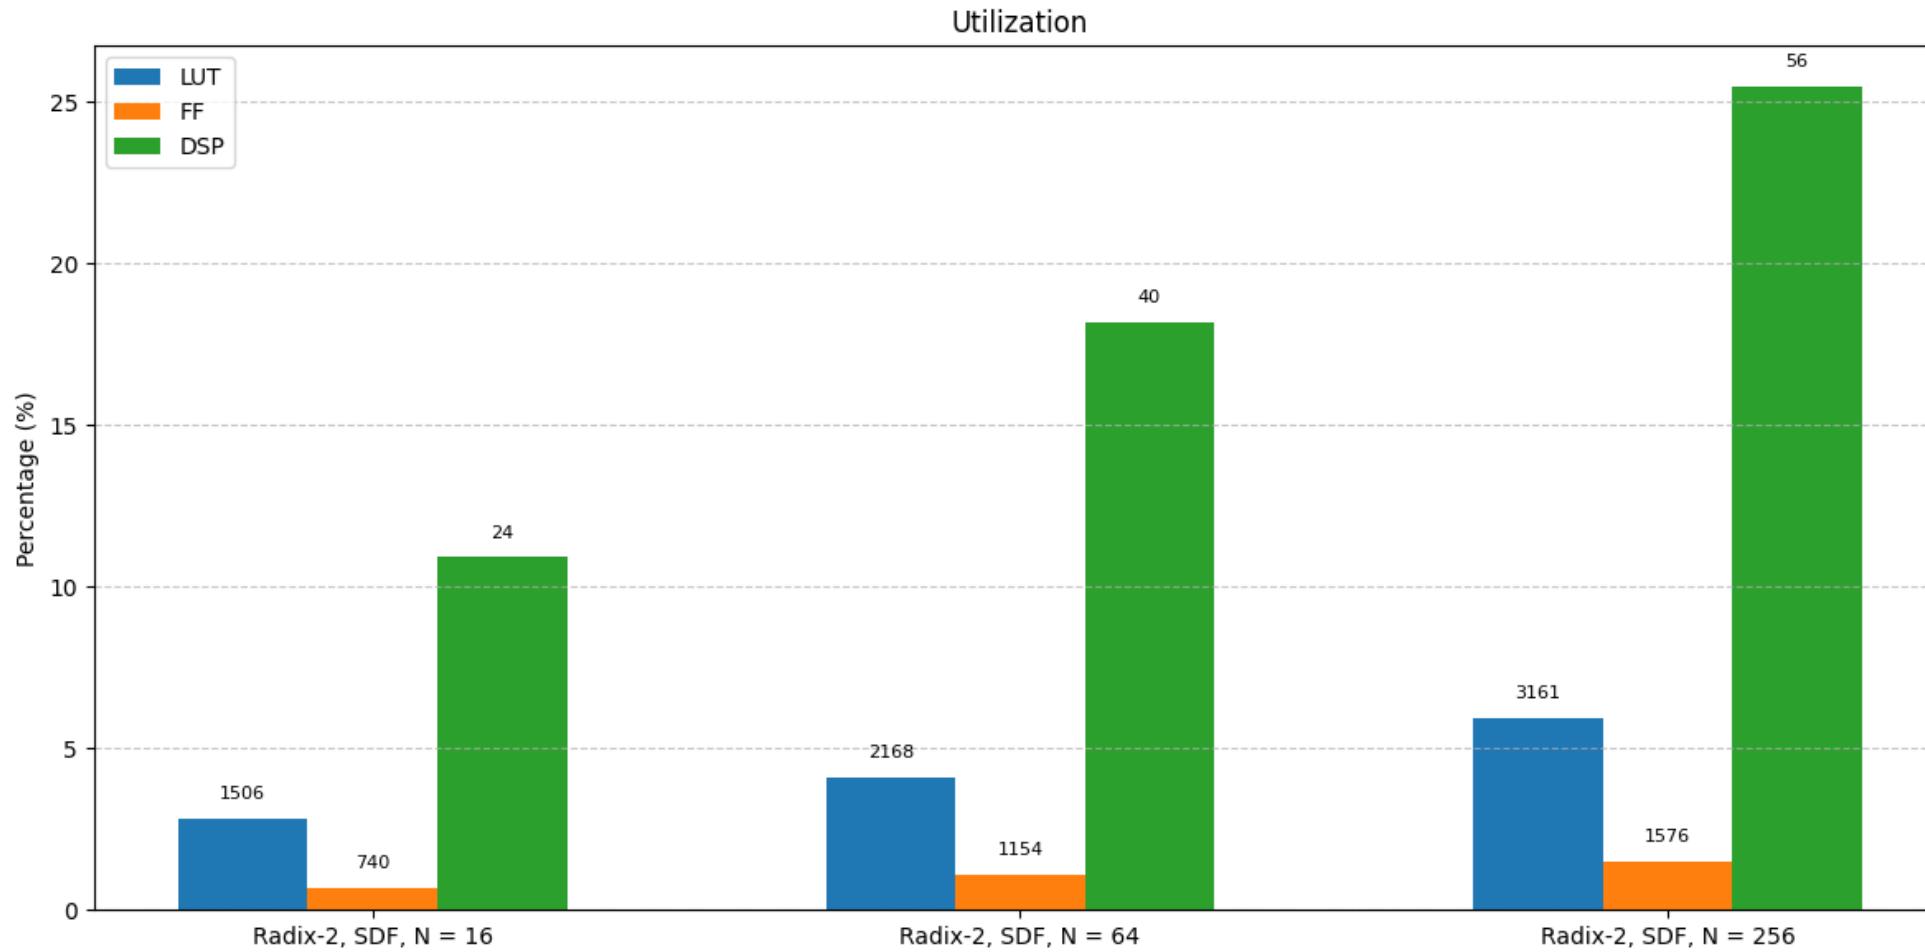


# NEW APPROACH

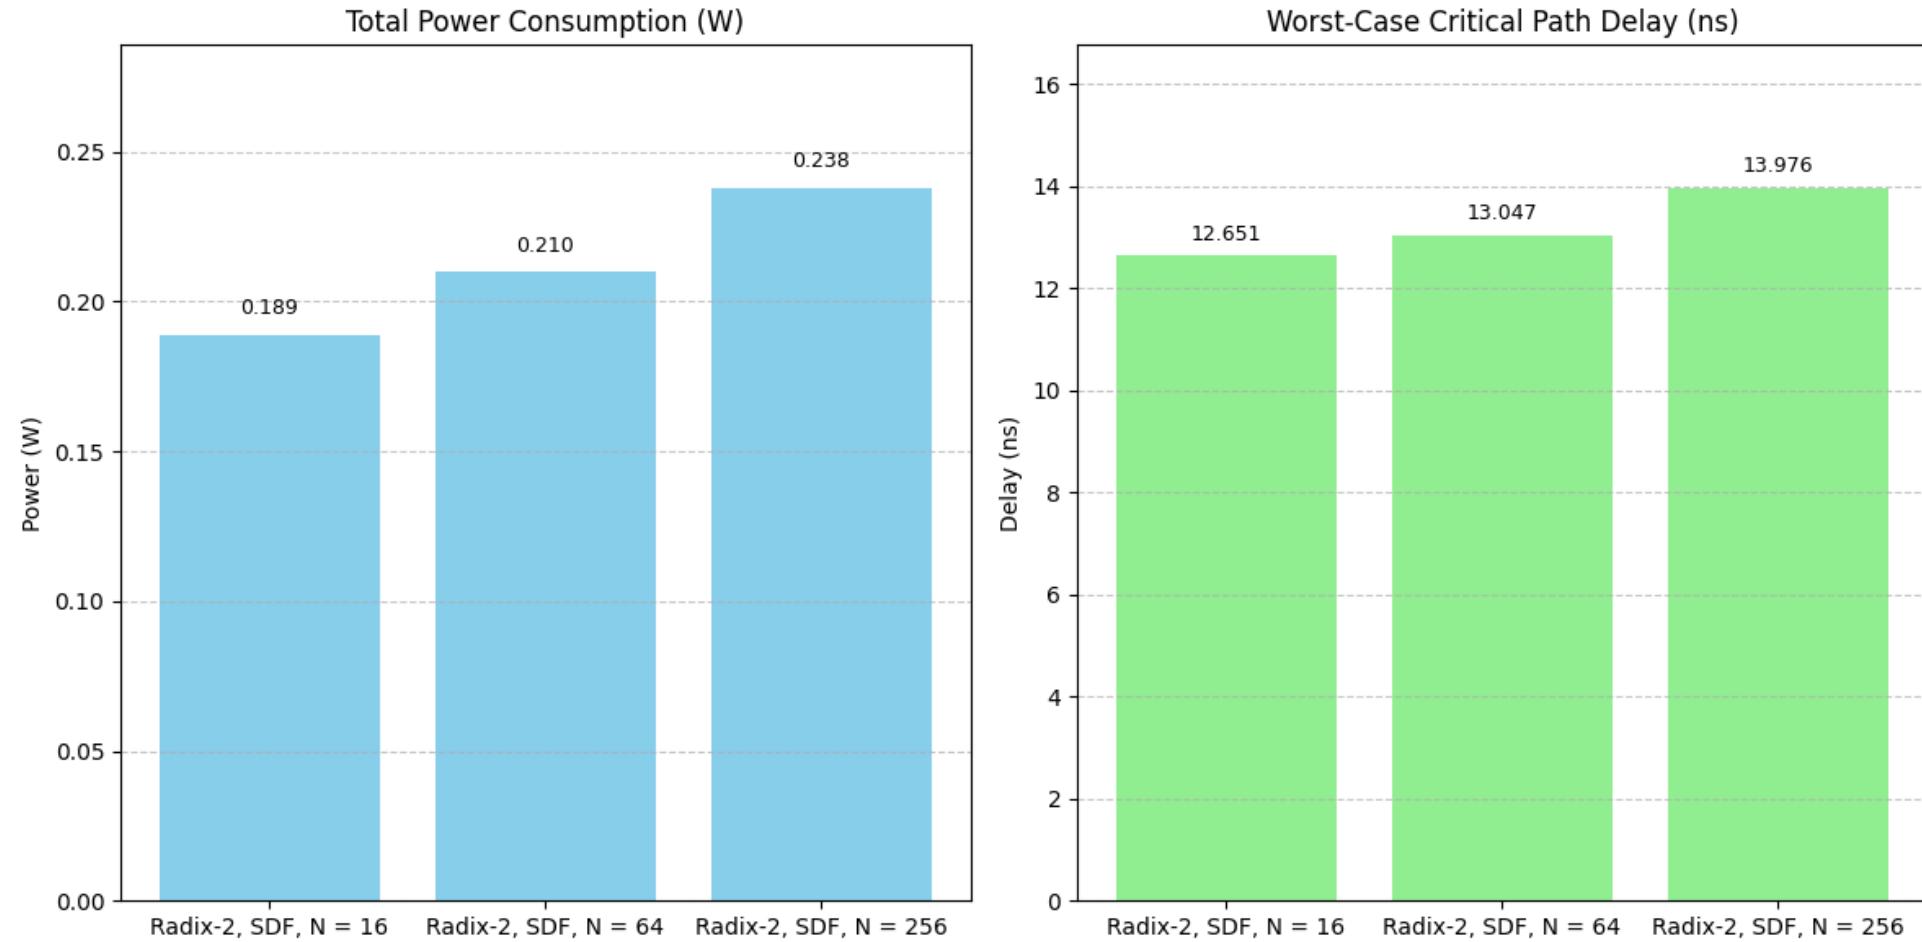
- Radix-2 SDF (Single path delay feedback)



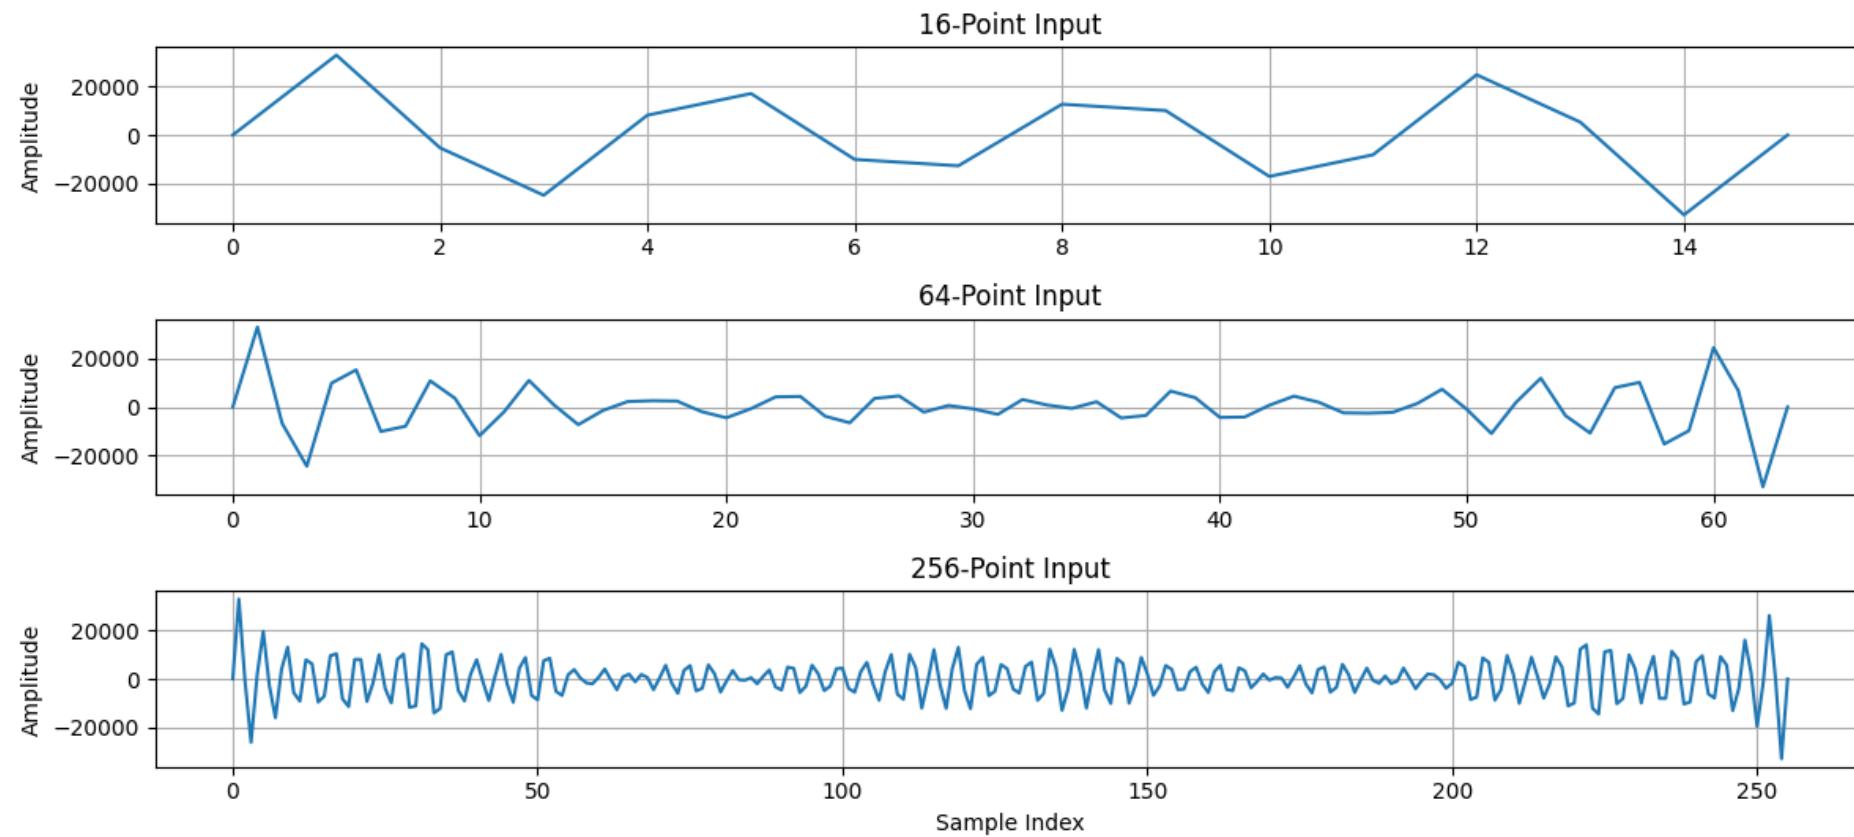
# RESULTS



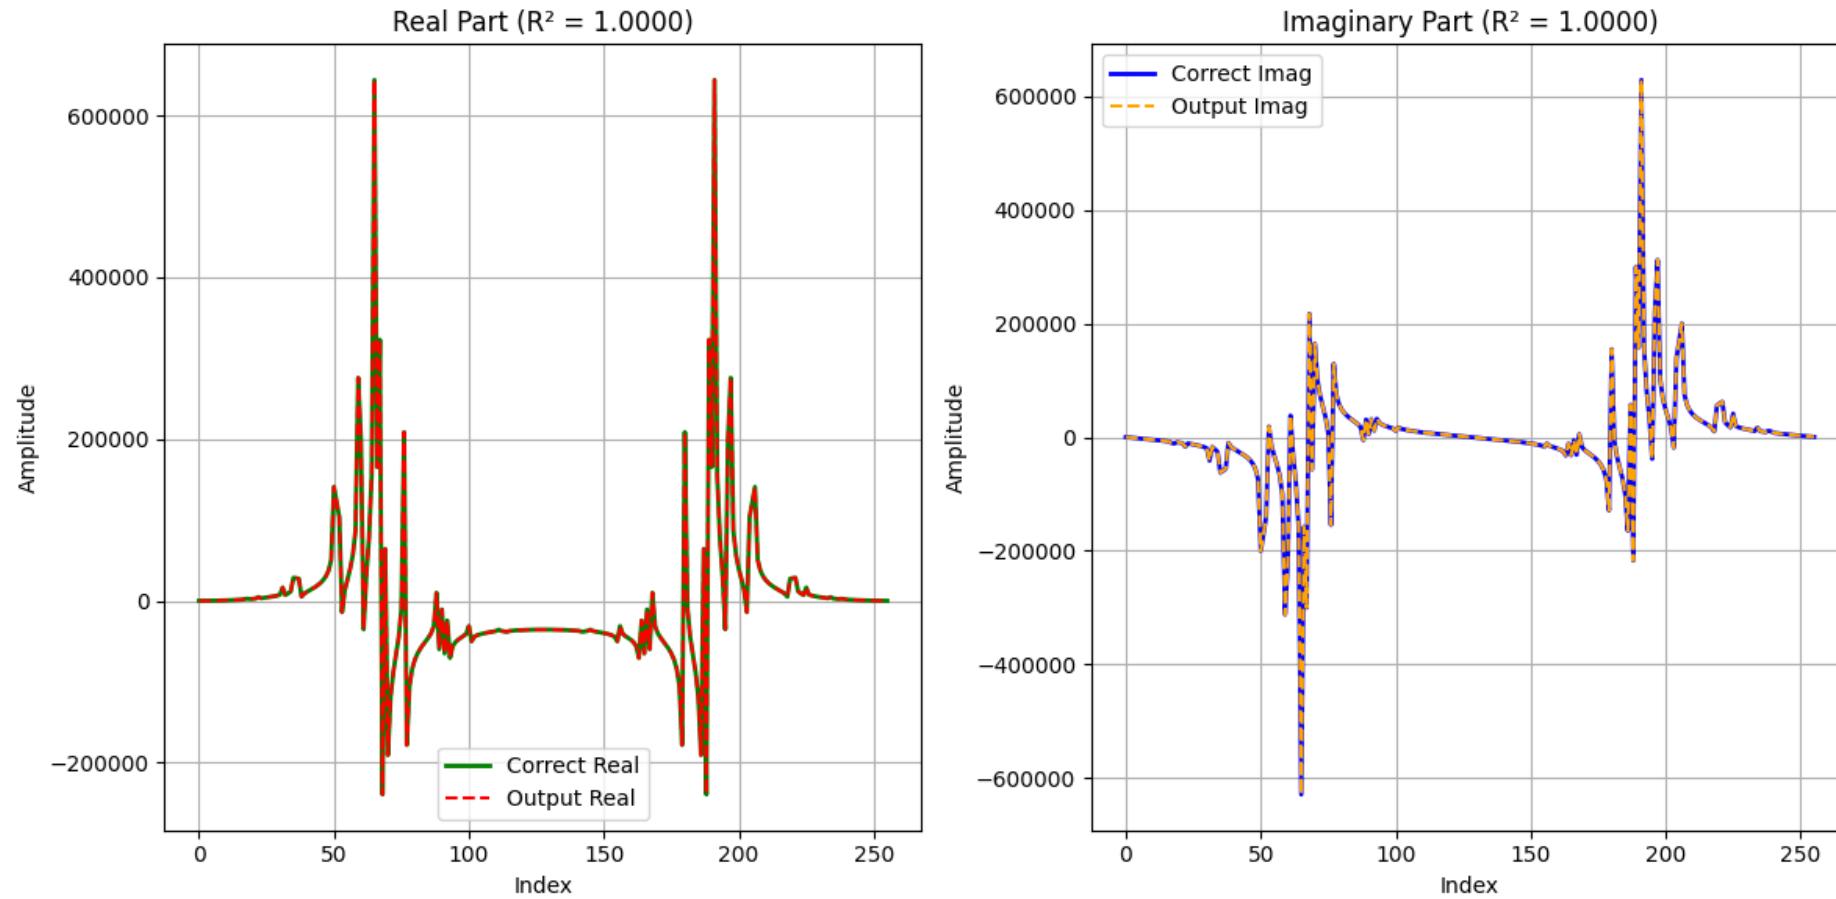
# RESULTS



# RESULTS



# RESULTS



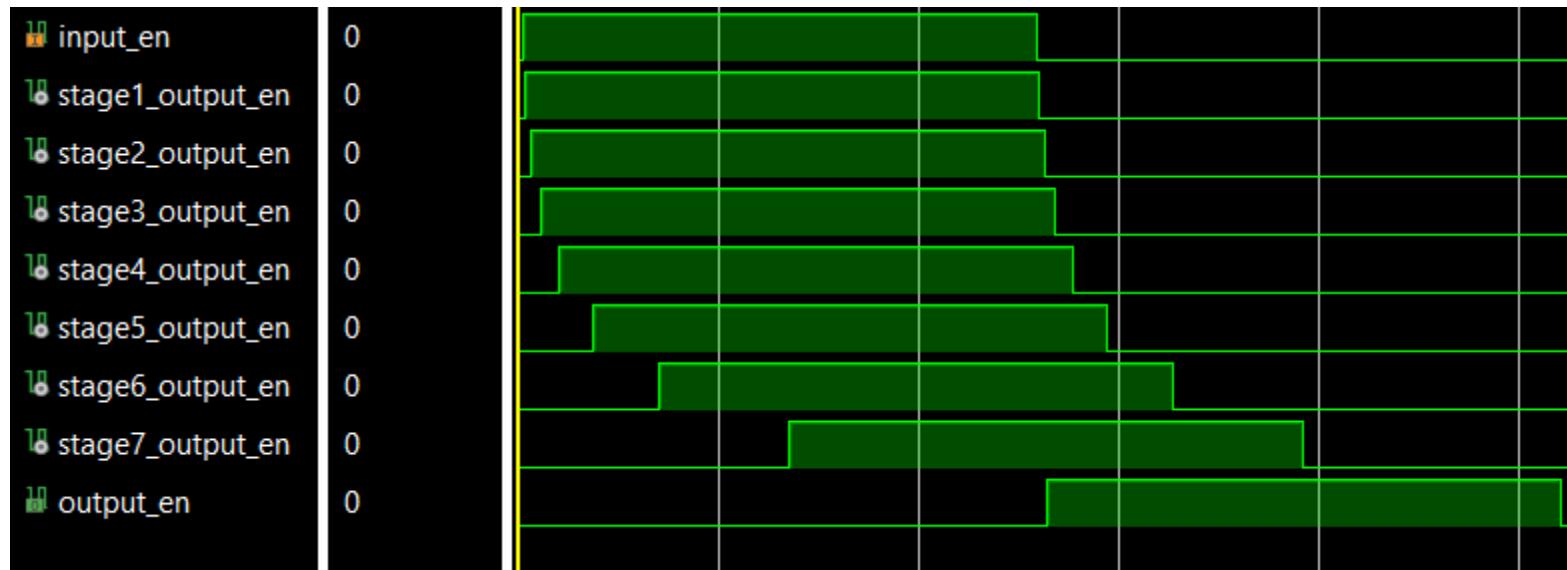
# EXECUTION TIME

- As you can see, larger delay buffers with each stage cause the latency from one to another to increase. For example, when measuring the offsets between output enables for consecutive stages we get  $1 + 3 + 5 + 9 + 17 + 31 + \dots + 256$  clock cycles ( $N = 256$ ). This can be generalized by the following formula:

$$N + 1 + \sum_{i=1}^{\log_2 N - 1} (2^i + 1)$$

which gets simplified to

$$2N - 2 + \log(N)$$

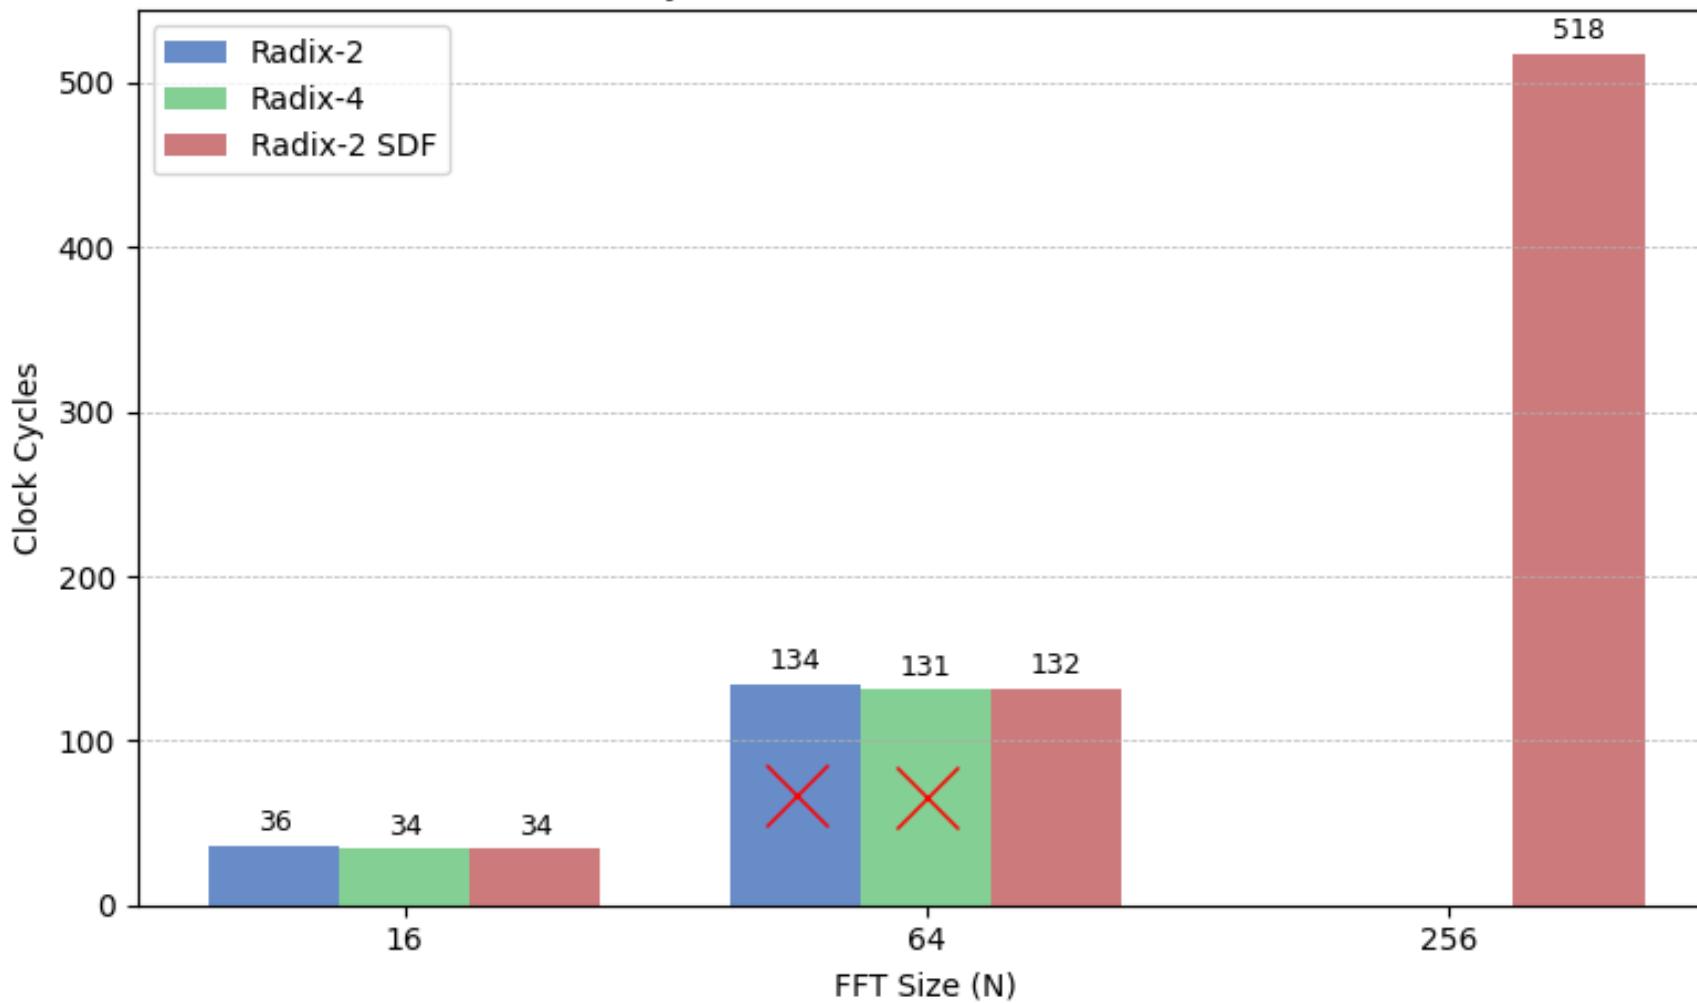


---

# THEORETICAL EXECUTION TIMES

- **Unrolled:**  $T = \log(N)$
  - **Unrolled with memory controller:**  $T = \frac{2N}{k} + \log_2(N)$   
( $k$  is the rate of transmission)
  - **SDF:**  $T = N + 1 + \sum_{i=1}^{\log_2 N - 1} (2^i + 1) = 2N - 2 + \log(N)$
-

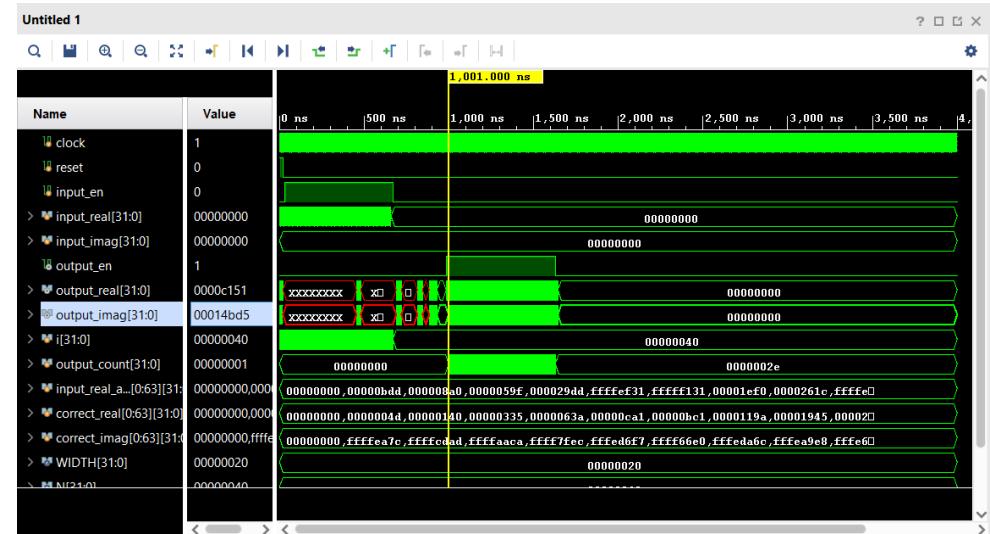
### Clock Cycles for Different FFT Architectures



# TOOLS USED

- We used Vivado for behavioral, post Implementation simulations
- The implementation was based on the ZedBoard FPGA
- All code is written with Verilog
- Algorithms were tested in software with C++
- Python for synthetic data (input, output, twiddles) and analysis (similarity scores, plots etc.)

Display Name	Preview	Vendor	File Ver...	Part
Nexys A7-100T		digilentinc.com	1.3	xc7a100tcs324-1
ZedBoard Zynq Evaluation and Development Kit Add Daughter Card <a href="#">Connections</a>		em.avnet.com	1.4	xc7z020clg484-1
Artix-7 AC701 Evaluation Platform Add Daughter Card <a href="#">Connections</a>		xilinx.com	1.4	xc7a200tfbg676-2
Kintex UltraScale+ KCU116 Evaluation Platform		xilinx.com	1.3	vcku116n-ffuh676-2-a



```
module fft_64_top #(  
    parameter WIDTH = 32,  
    parameter N = 64  
)  
(  
    input          clock,  
    input          reset,  
    input          input_en,  
    input [WIDTH-1:0] input_real,  
    input [WIDTH-1:0] input_imag,  
    output         output_en,  
    output [WIDTH-1:0] output_real,  
    output [WIDTH-1:0] output_imag  
  
    // Internal signals between stages  
    wire stage1_output_en;  
    wire [WIDTH-1:0] stage1_output_real;  
    wire [WIDTH-1:0] stage1_output_imag;  
  
    wire stage2_output_en;  
    wire [WIDTH-1:0] stage2_output_real;  
    wire [WIDTH-1:0] stage2_output_imag;  
);
```

---

# PROBLEMS FACED

- Memory bandwidth bottlenecks in the unrolled versions forced us to implement a memory controller to feed the data. This decision changed the algorithm's complexity from  $O(\log N)$  to  $O(N)$ . Theoretically, using an AXI master interface could improve the data rate (bigger  $k$ ), but even then, the unrolled algorithm remains not scalable for small FPGAs
- We encountered difficulties implementing the SDF algorithm for the Radix-4 FFT.

